

Learning list functions through program induction

Joshua Rule

Learning as Program Induction Workshop
CogSci 2018, 25 July 2018



Computational
Cognitive Science Group

 brain+cognitive
sciences



CENTER FOR
Brains
Minds +
Machines



This talk

- ▶ learning as programming
- ▶ bootstrapping the LOT with term rewriting
- ▶ toward a model of conceptual change

This talk

- ▶ learning as programming
- ▶ bootstrapping the LOT with term rewriting
- ▶ toward a model of conceptual change



Build causal theories from sparse evidence

Navigate complex environments

Recognize objects, reason cross-modally

Tie shoes, make bed, set table

Introspect on beliefs and desires

whisper, shout, sing, joke

Build towers, sandcastles, & Lego cars

Use light switches, door knobs, & smartphones

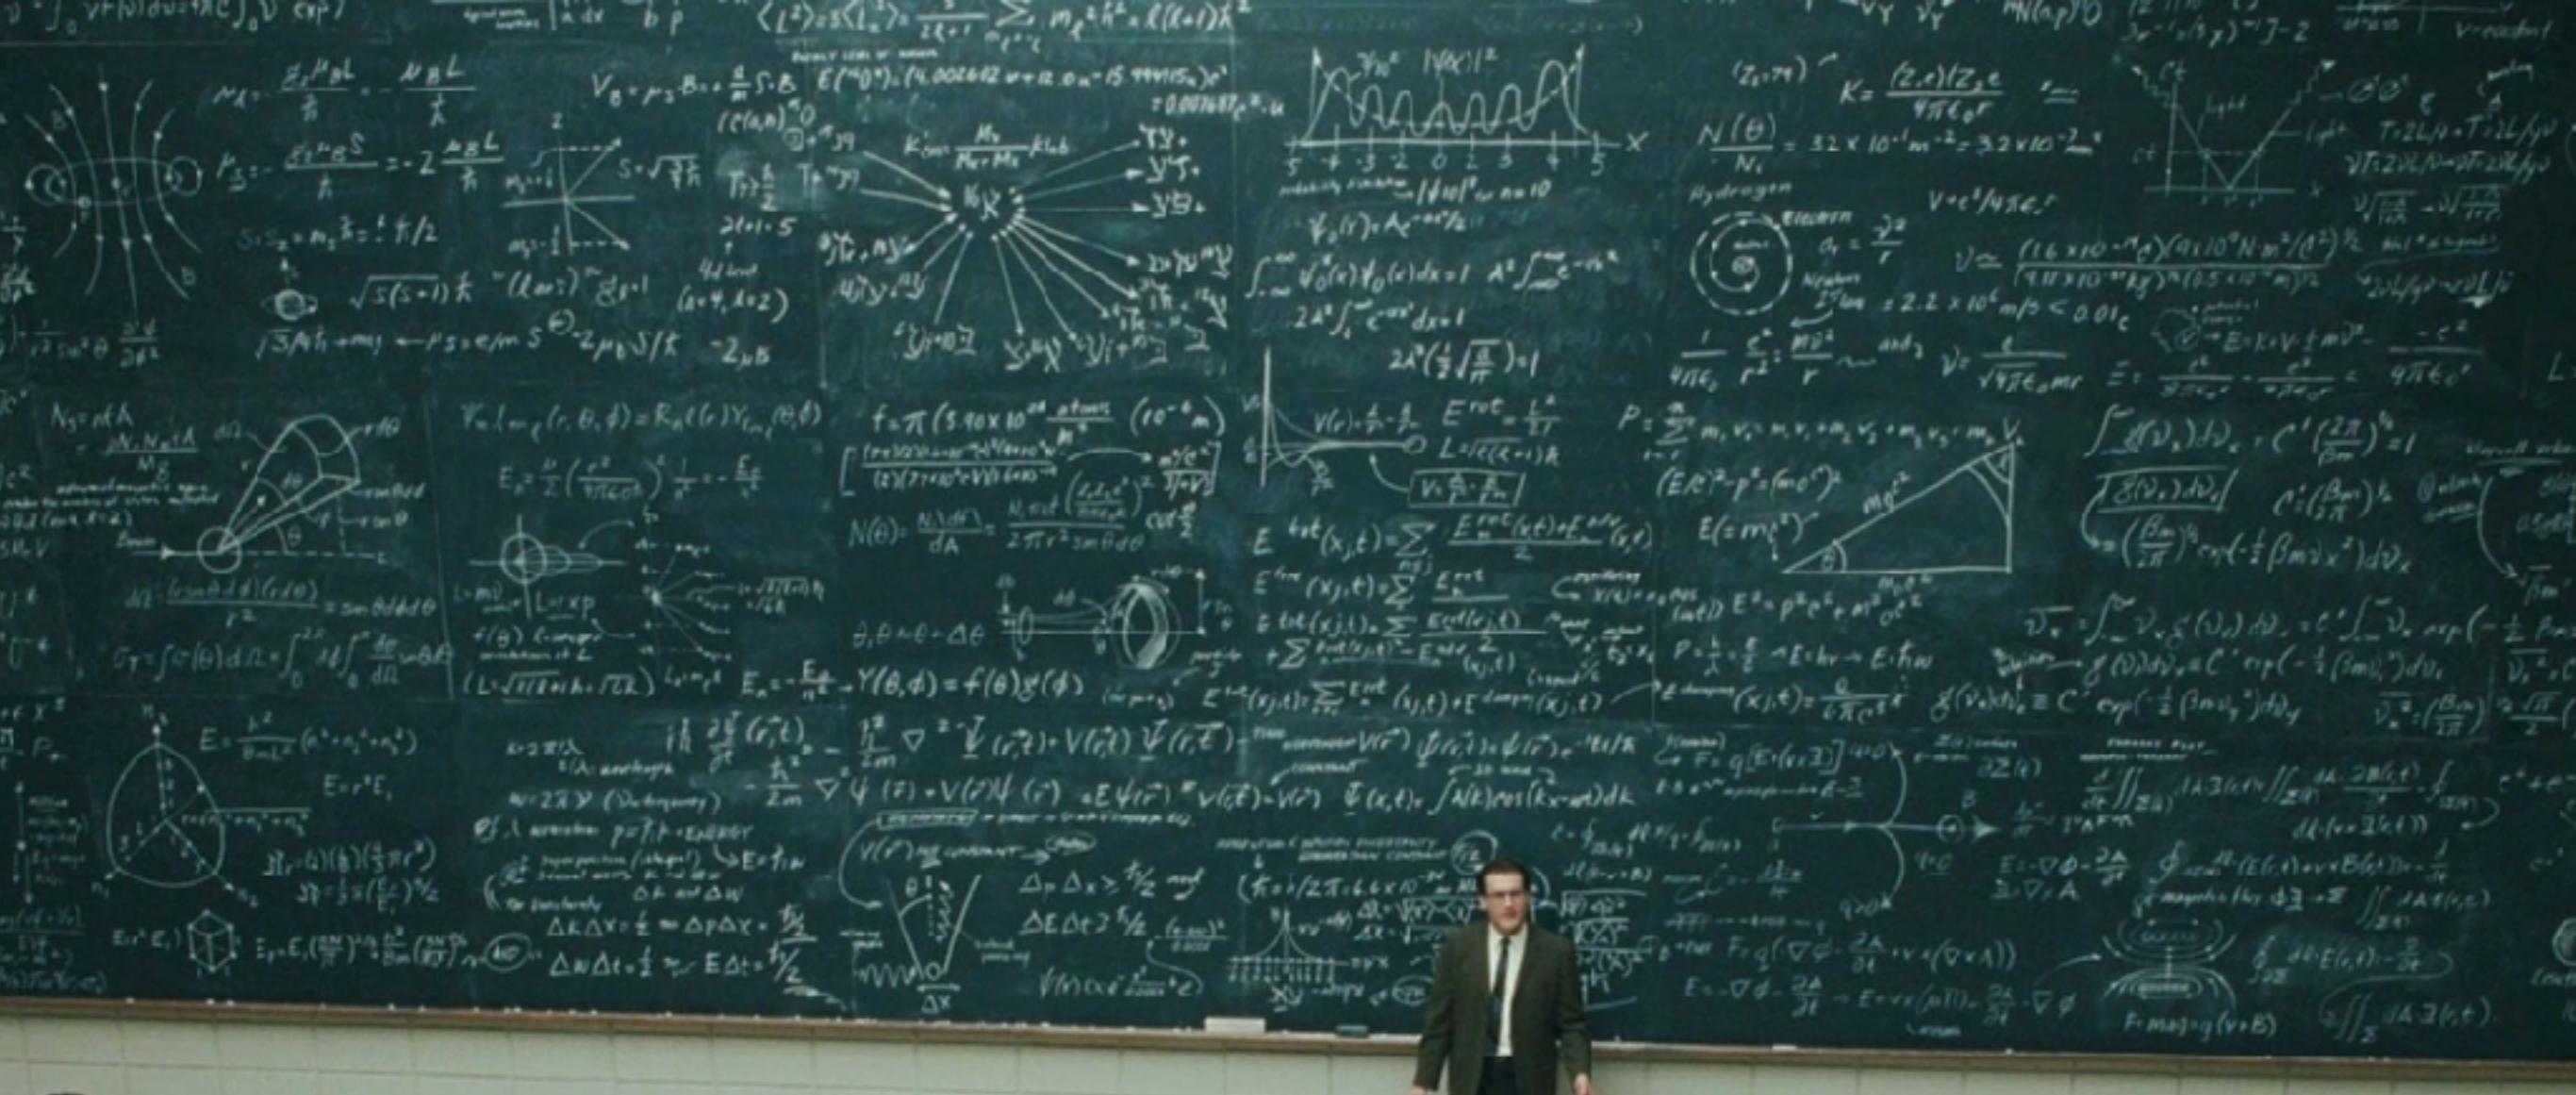
Talk about dinosaurs, trucks, and fairy tales

Play with others, share, determine ownership

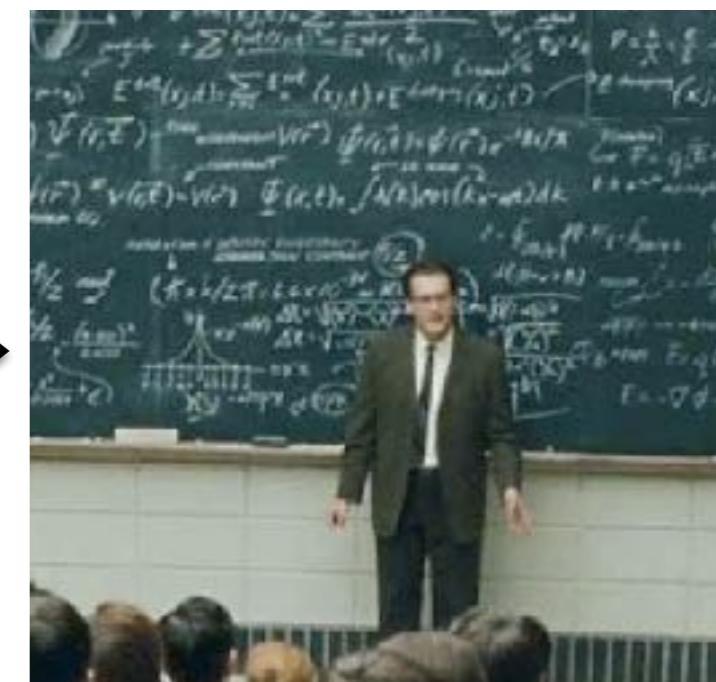
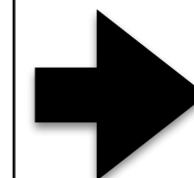
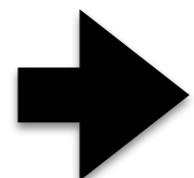
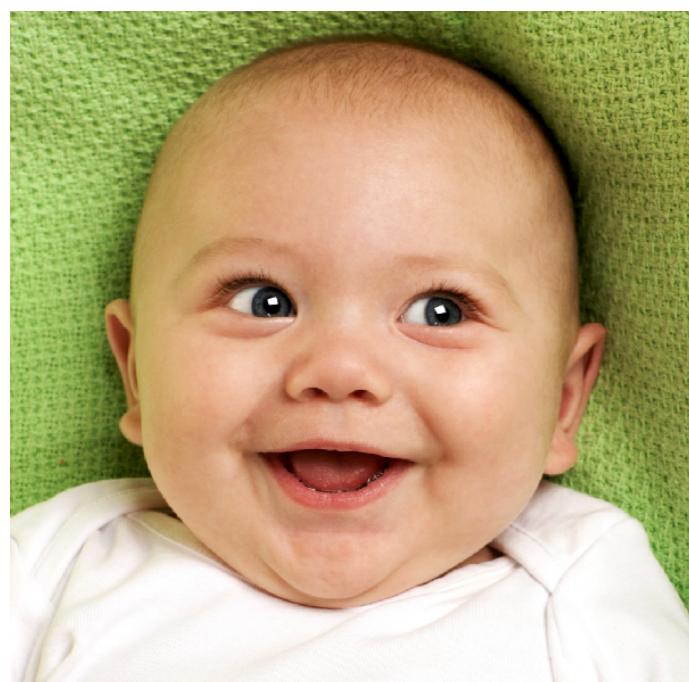
Walk, run, skip, dance, somersault

Use natural language





The Puzzle of Learning & Cognitive Development

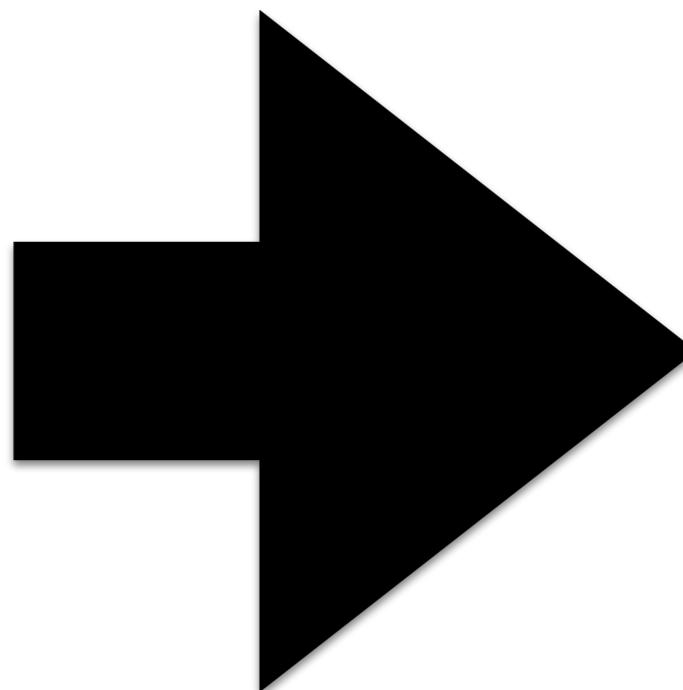


(see Tenenbaum, Kemp, Griffiths, Goodman, 2011; Carey, 2009)

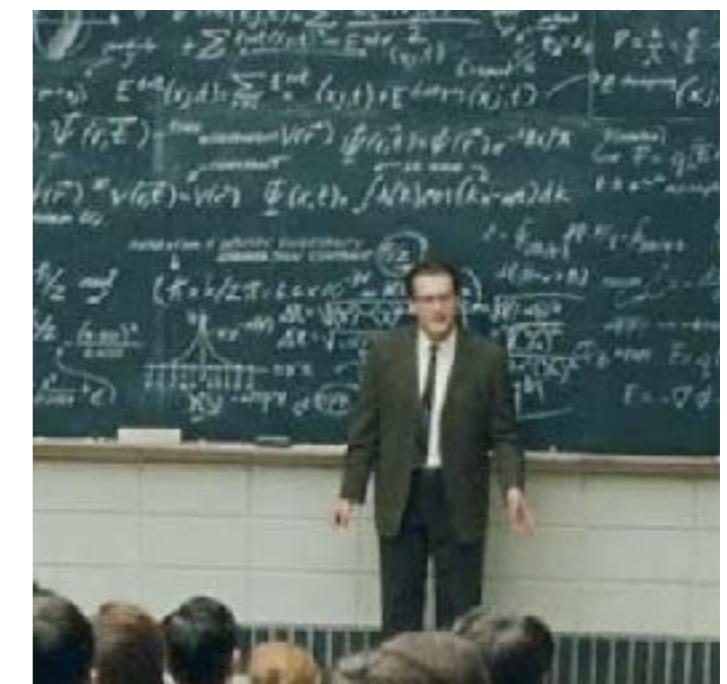
The Puzzle of Learning & Cognitive Development



Initial State



Learning Mechanism



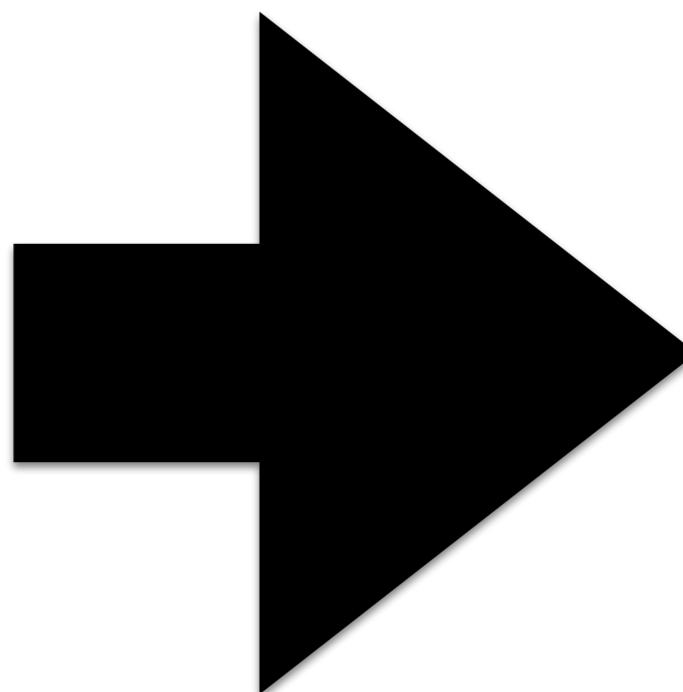
Final State

(see Tenenbaum, Kemp, Griffiths, Goodman, 2011; Carey, 2009)

The Puzzle of Learning & Cognitive Development



Initial State



Learning Mechanism



Final State

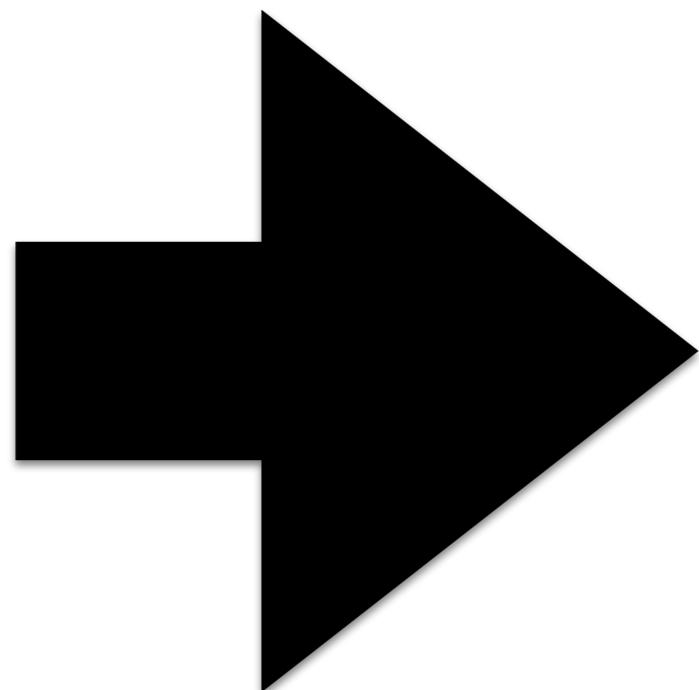
(see Tenenbaum, Kemp, Griffiths, Goodman, 2011; Carey, 2009)

The Puzzle of Learning & Cognitive Development



Initial State

LOT?



Learning Mechanism

***inductive
bootstrapping?***



Final State

LOT'?

(Carey, 1985, 2009; Carey, Spelke, 1994)

(Fodor, 1975; Turing, 1936; Fodor & Pylyshyn, 1988; Goodman, Tenenbaum, & Gerstenberg, 2015)

Three questions about learning in the LOT

1. How are concepts represented?
2. How are changes proposed?
3. How are proposals assessed?

Three questions about learning in the LOT

1. How are concepts represented?

- **programs in some language**

2. How are changes proposed?

3. How are proposals assessed?

(Goodman, Tenenbaum, Feldman, & Griffiths, 2008; Piantadosi, Tenenbaum, & Goodman, 2012, 2016; Lake, Salakhutdinov, & Tenenbaum, 2015; Kemp & Tenenbaum, 2008; Dechter, Malmaud, Adams, & Tenenbaum, 2013; Rule, Dechter, & Tenenbaum, 2015; Piantadosi, unpub.; Ullman, Goodman, & Tenenbaum, 2012; Goodman, Ullman, & Tenenbaum, 2011)

Three questions about learning in the LOT

1. How are concepts represented?

- programs in some language

2. How are changes proposed?

- **small, random syntactic changes to a concept definition**

3. How are proposals assessed?

(Goodman, Tenenbaum, Feldman, & Griffiths, 2008; Piantadosi, Tenenbaum, & Goodman, 2012, 2016; Lake, Salakhutdinov, & Tenenbaum, 2015; Kemp & Tenenbaum, 2008; Dechter, Malmaud, Adams, & Tenenbaum, 2013; Rule, Dechter, & Tenenbaum, 2015; Piantadosi, unpub.; Ullman, Goodman, & Tenenbaum, 2012; Goodman, Ullman, & Tenenbaum, 2011)

Three questions about learning in the LOT

1. How are concepts represented?

- programs in some language

2. How are changes proposed?

- small, random syntactic changes to a concept definition

3. How are proposals assessed?

- **accuracy & description length (& sometimes efficiency)**

(Goodman, Tenenbaum, Feldman, & Griffiths, 2008; Piantadosi, Tenenbaum, & Goodman, 2012, 2016; Lake, Salakhutdinov, & Tenenbaum, 2015; Kemp & Tenenbaum, 2008; Dechter, Malmaud, Adams, & Tenenbaum, 2013; Rule, Dechter, & Tenenbaum, 2015; Piantadosi, unpub.; Ullman, Goodman, & Tenenbaum, 2012; Goodman, Ullman, & Tenenbaum, 2011)

Three questions about learning in the LOT

1. How are concepts represented?

- programs in some language

2. How are changes proposed?

- small, random syntactic changes to a concept definition

3. How are proposals assessed?

- accuracy & description length (& sometimes efficiency)

(Goodman, Tenenbaum, Feldman, & Griffiths, 2008; Piantadosi, Tenenbaum, & Goodman, 2012, 2016; Lake, Salakhutdinov, & Tenenbaum, 2015; Kemp & Tenenbaum, 2008; Dechter, Malmaud, Adams, & Tenenbaum, 2013; Rule, Dechter, & Tenenbaum, 2015; Piantadosi, unpub.; Ullman, Goodman, & Tenenbaum, 2012; Goodman, Ullman, & Tenenbaum, 2011)

This talk

- ▶ learning as programming
- ▶ bootstrapping the LOT with term rewriting
- ▶ toward a model of conceptual change

A nagging problem

LOT + bootstrapping = LOT'

A nagging problem

$$LOT + bootstrapping = LOT'$$

modeled as

$$\begin{matrix} \textit{Prog. Lang.} \\ \& \\ \textit{Library} \end{matrix} + \begin{matrix} \textit{stochastic} \\ \textit{search} \end{matrix} = \begin{matrix} \textit{Prog. Lang.} \\ \& \\ \textit{Library}' \end{matrix}$$

Three questions about learning in the LOT

1. How are concepts represented?

- **programs in some language**

2. How are changes proposed?

- small, random syntactic changes to a concept definition

3. How are proposals assessed?

- accuracy & description length (& sometimes efficiency)

(Goodman, Tenenbaum, Feldman, & Griffiths, 2008; Piantadosi, Tenenbaum, & Goodman, 2012, 2016; Lake, Salakhutdinov, & Tenenbaum, 2015; Kemp & Tenenbaum, 2008; Dechter, Malmaud, Adams, & Tenenbaum, 2013; Rule, Dechter, & Tenenbaum, 2015; Piantadosi, unpub.; Ullman, Goodman, & Tenenbaum, 2012; Goodman, Ullman, & Tenenbaum, 2011)

Three questions about learning in the LOT

1. How are concepts represented?

- **programs in some *fixed* language**

2. How are changes proposed?

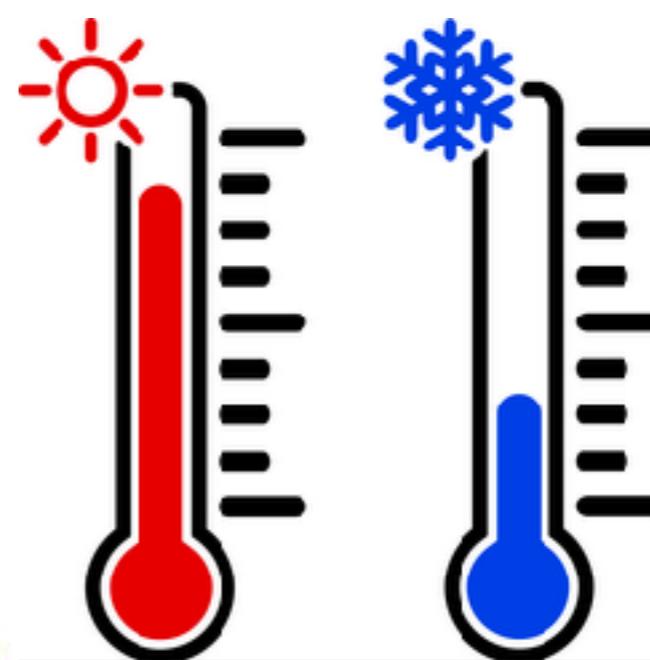
- small, random syntactic changes to a concept definition

3. How are proposals assessed?

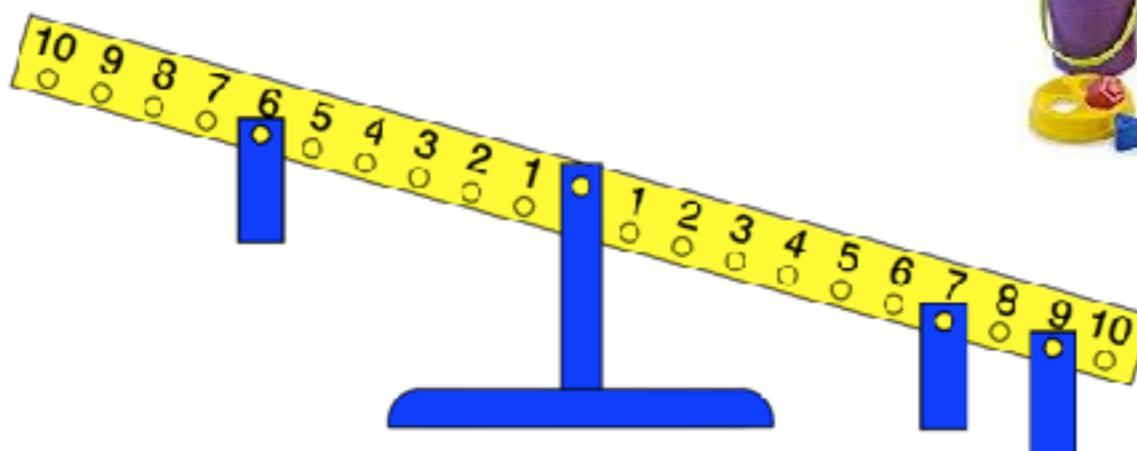
- accuracy & description length (& sometimes efficiency)

(Goodman, Tenenbaum, Feldman, & Griffiths, 2008; Piantadosi, Tenenbaum, & Goodman, 2012, 2016; Lake, Salakhutdinov, & Tenenbaum, 2015; Kemp & Tenenbaum, 2008; Dechter, Malmaud, Adams, & Tenenbaum, 2013; Rule, Dechter, & Tenenbaum, 2015; Piantadosi, unpub.; Ullman, Goodman, & Tenenbaum, 2012; Goodman, Ullman, & Tenenbaum, 2011)

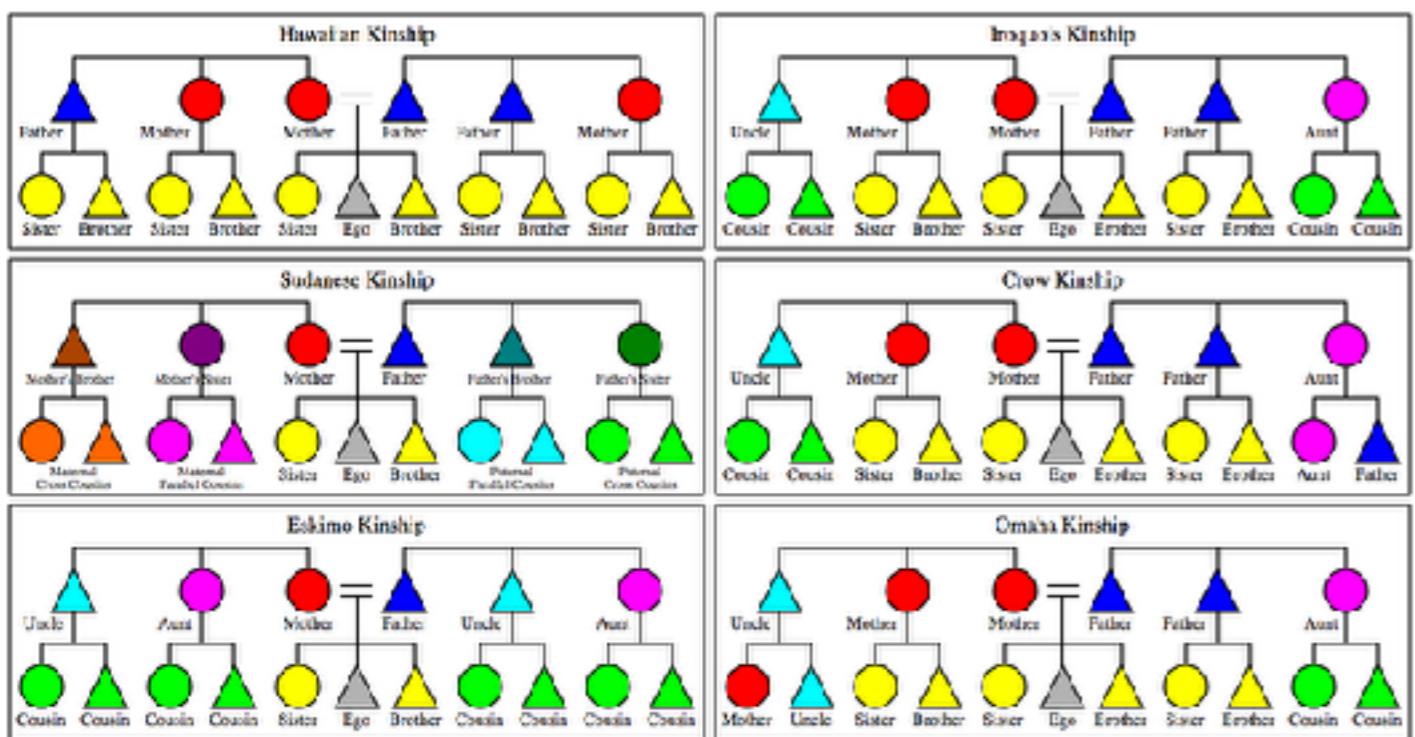
	pollen ♂	
	B	b
pistil ♀	B	BB
b	Bb	bb



AABABA
ABAAAB
AABAAB
ABAAAB
ABAABA



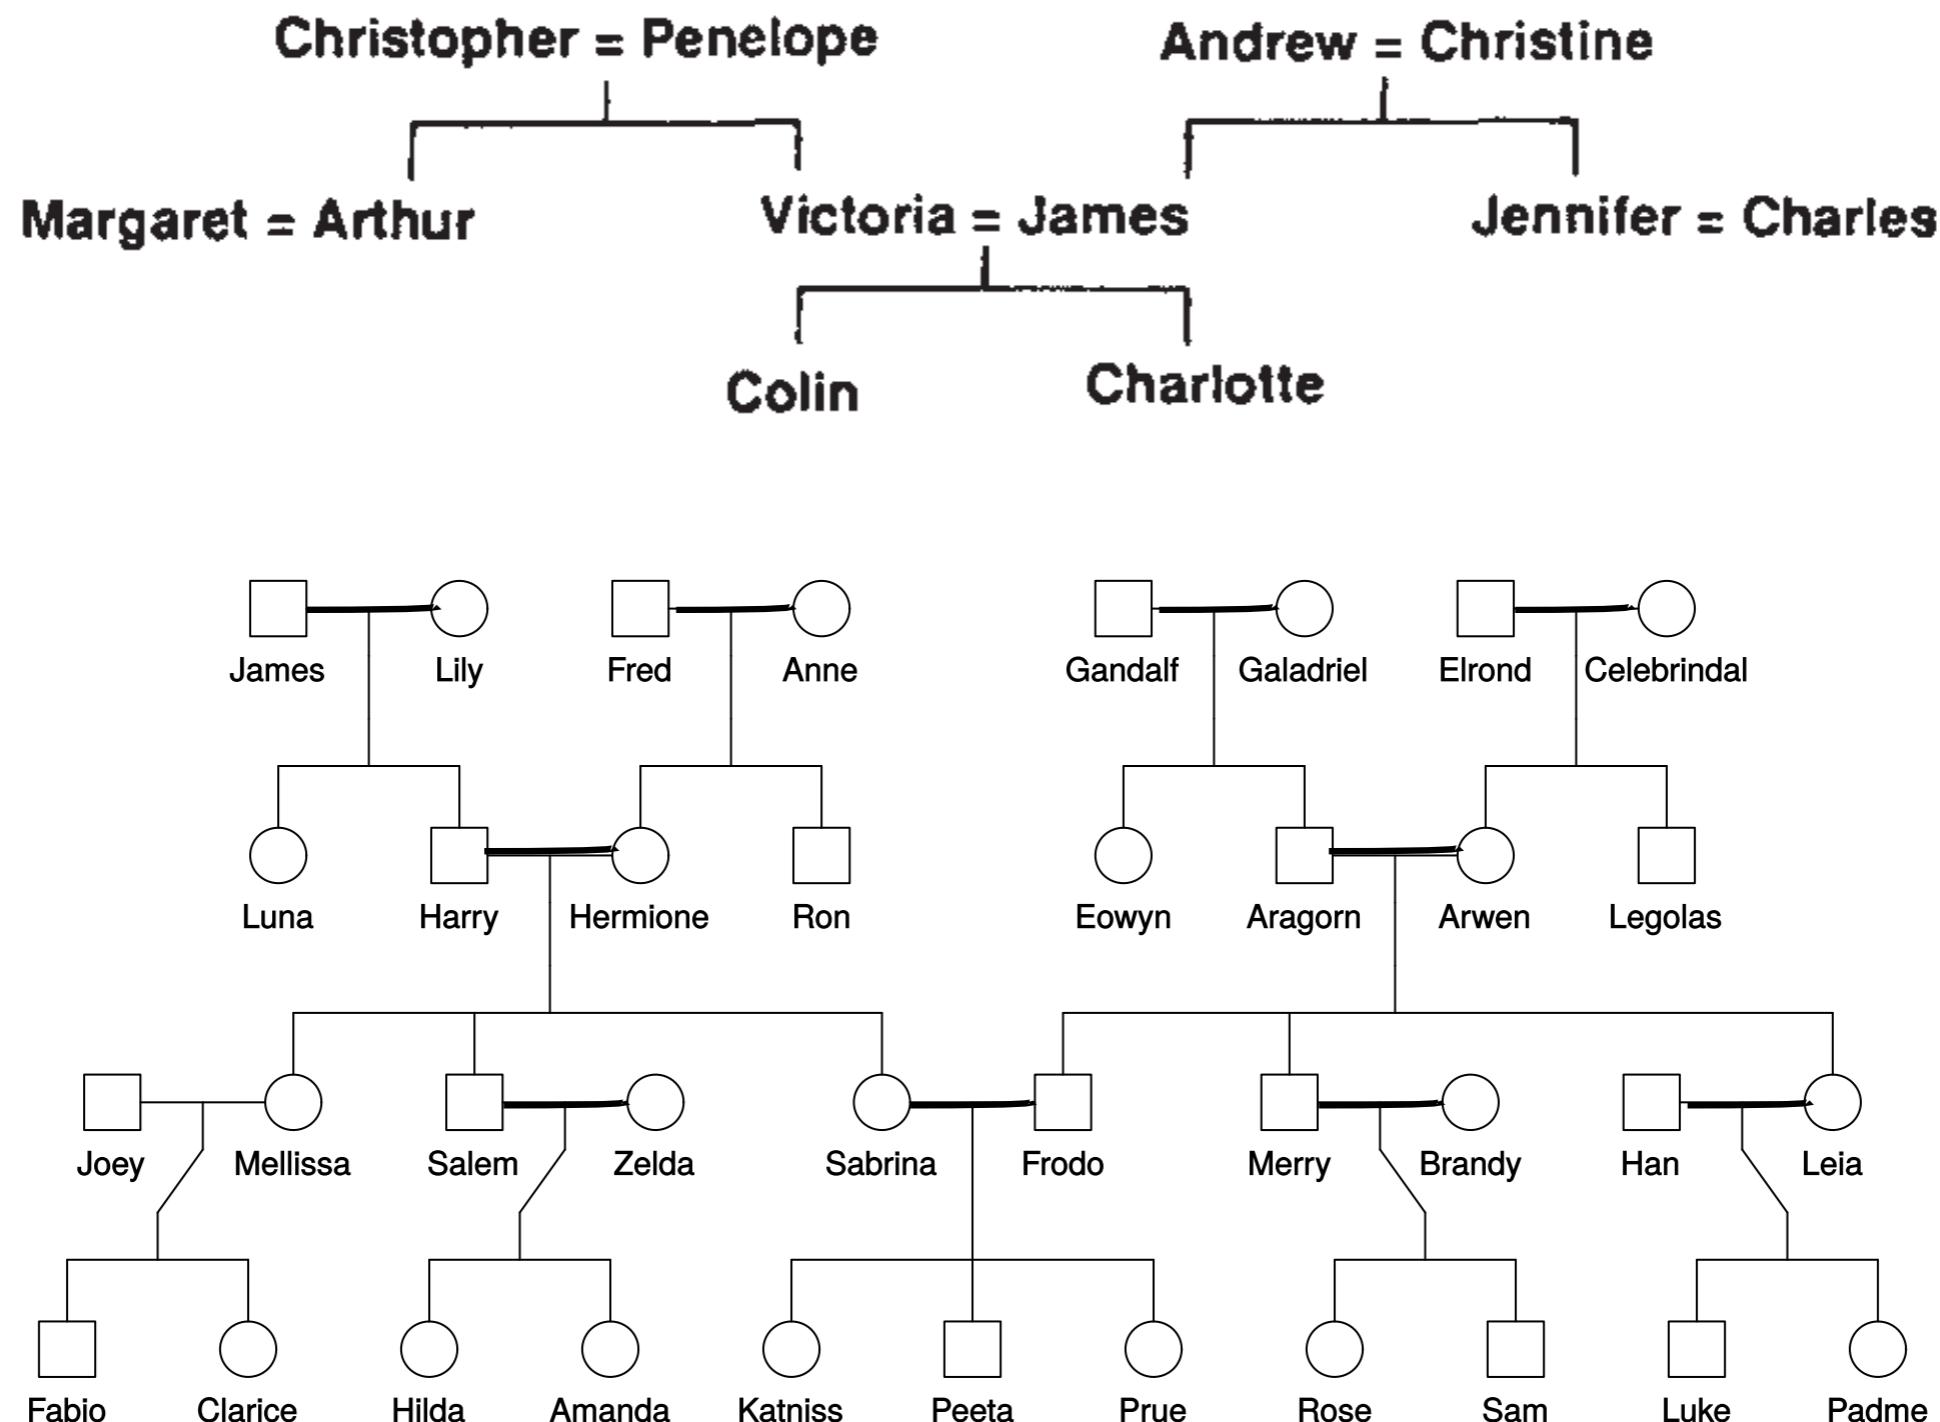
six-hundred-forty-seven-thousand-nine-hundred-sixteen



Kinship is a great space for studying conceptual change

Definite Gender	boy, girl, man, woman
Generic Gender	male, female
Definite Nuclear	brother, sister, mother, father, husband, wife, son, daughter
Generic Nuclear	sibling, spouse, parent, child
Definite Extended	aunt, uncle, nephew, niece, grandmother, grandfather, granddaughter, grandson, grandnephew, grandniece
Generic Extended	grandparent, grandchild, cousin
Structurally Recursive	great-aunt, great-uncle, great-grandfather, great-grandmother, great-grandparent, great-granddaughter, great-grandson, great-grandchild, great-great-, great-great-great, ...
Linearly Recursive	ancestor, descendant
Nonlinearly Recursive	relative, blood relative, in-law, m^{th} cousin n^{th} removed, step-relations

typical kinship data



potential kinship data

:

true → husband(Christopher, Penelope)

true → cousin(Rose, Luke)

true → uncle(Arthur, Colin)

true → brother(Arthur, Victoria)

true → man(Arthur)

true → girl(Charlotte)

true → dad(Joey, Clarice)

true → brother(Sam, Rose)

true → great-uncle(Ron, Katniss)

true → sister(Katniss, Prue)

true → sister(Prue, Katniss)

true → husband(James, Victoria)

true → sister(Rose, Sam)

false → sister(Sam, Rose)

:

potential kinship grammar

potential kinship grammar

male

female

spouse

parent

potential kinship grammar

male(Aragorn)

female(Arwen)

spouse(Aragorn, Arwen)

parent(Elrond, Arwen)

potential kinship grammar

male(Aragorn)

female(Arwen)

spouse(Aragorn, Arwen)

parent(Elrond, Arwen)

:

and(male(x), spouse(x, y)) → husband(x, y)

and(female(x), spouse(x, y)) → wife(y, y)

and(female(x), sibling(x, y)) → sister(x, y)

and(male(x), sibling(x, y)) → brother(x, y)

and(male(x), parent(x, y)) → father(x, y)

and(female(x), parent(x, y)) → mother(x, y)

and(male(x), parent(y, x)) → son(x, y)

and(female(x), parent(y, x)) → daughter(x, y)

and(parent(z,y), parent(z,x)) → sibling(x, y)

:

parent(x, y) → ancestor(x, y)

and(parent(x, y), ancestor(y, z)) → ancestor(x, y)

:

and(ancestor(x, y), ancestor(x, z)) → blood_relative(y, z)

Three questions about learning in the LOT

1. How are concepts represented?

- **programs in some *fixed* language**

2. How are changes proposed?

- small, random syntactic changes to a concept definition

3. How are proposals assessed?

- accuracy & description length (& sometimes efficiency)

(Goodman, Tenenbaum, Feldman, & Griffiths, 2008; Piantadosi, Tenenbaum, & Goodman, 2012, 2016; Lake, Salakhutdinov, & Tenenbaum, 2015; Kemp & Tenenbaum, 2008; Dechter, Malmaud, Adams, & Tenenbaum, 2013; Rule, Dechter, & Tenenbaum, 2015; Piantadosi, unpub.; Ullman, Goodman, & Tenenbaum, 2012; Goodman, Ullman, & Tenenbaum, 2011)

Three questions about learning in the LOT

1. How are concepts represented?

- **programs in some *adaptive* language**

2. How are changes proposed?

- small, random syntactic changes to a concept definition

3. How are proposals assessed?

- accuracy & description length (& sometimes efficiency)

(Goodman, Tenenbaum, Feldman, & Griffiths, 2008; Piantadosi, Tenenbaum, & Goodman, 2012, 2016; Lake, Salakhutdinov, & Tenenbaum, 2015; Kemp & Tenenbaum, 2008; Dechter, Malmaud, Adams, & Tenenbaum, 2013; Rule, Dechter, & Tenenbaum, 2015; Piantadosi, unpub.; Ullman, Goodman, & Tenenbaum, 2012; Goodman, Ullman, & Tenenbaum, 2011)

Term Rewriting Systems

$$TRS = (\Sigma, R)$$

Term Rewriting Systems

Signature:

- a set of primitives
- what things exist
- syntax

$$TRS = (\Sigma, R)$$


Term Rewriting Systems

Signature:

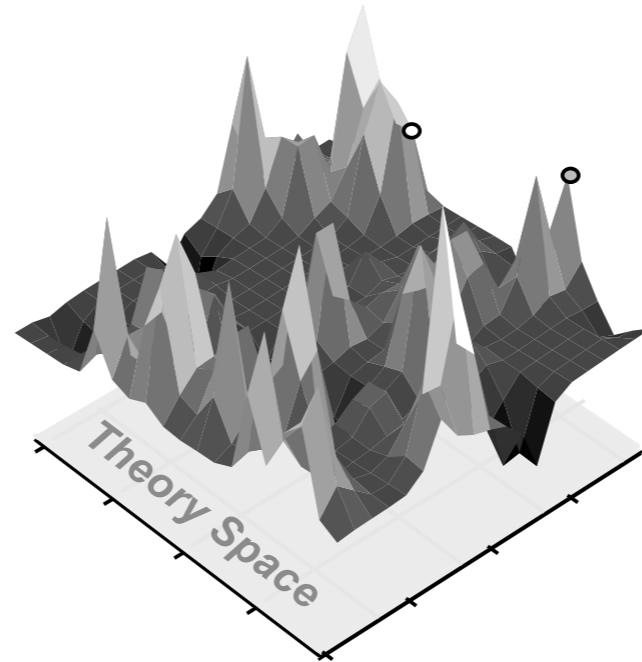
- a set of primitives
- what things exist
- syntax

$$TRS = (\Sigma, R)$$

Rules:

- a list of rewrite rules
- how things behave
- semantics

Stochastic search over TRSs



- remove a symbol s from Σ_{i-1} and all rules involving s from R_{i-1}
- add a symbol s to Σ_i
- generate a new rule r and add it to R_i
- remove a rule r from R_{i-1}

One solution: models LOTs as Term Rewriting Systems (TRSs)

$$LOT + \text{bootstrapping} = LOT'$$

modeled as

$$TRS + \text{stochastic search} = TRS'$$

This talk

- ▶ learning as programming
- ▶ bootstrapping the LOT with term rewriting
- ▶ toward a model of conceptual change



Josh Tenenbaum



Human learning is astonishingly efficient, mastering complex systems of facts and rules using surprisingly few trials (Tenenbaum, Kemp, Griffiths, & Goodman, 2011). It seems to do so by compositionally recombining smaller parts (e.g., Schulz, Tenenbaum, Duvenaud, Speekenbrink, & Tenenbaum, 2017) and incrementally developing new representations (Griffiths, & Lagnado, 2013). One way to explain these abilities is to model the concept of learning as a program-like structure that takes observations, also known as program induction (Dechter, Malmaud, Adams, & Tenenbaum, 2013; Lake, Salakhutdinov, & Tenenbaum, 2015; Piantadosi, Tenenbaum, & Goodman, 2016). Program induction algorithms have been used to model unsupervised learning and sequence learning (Ellis, Dechter, & Tenenbaum, 2015; Romano, Salles, Amalric, Dehaene, Sigman, & Figueria, 2017), to support one-shot inferences (Lake et al., 2015), and to investigate the primitives of thought (Piantadosi et al., 2016).



Steve Piantadosi



um, Feldman, Gershman, and Goldstone (2017) have shown that people can learn to use such semantic relations to make predictions about novel situations. However, it is not clear whether these findings generalize to other domains or the preexisting knowledge that people have. In this paper, we show that people can indeed learn to use semantic relations to make predictions about novel situations, even when they have no prior knowledge of the relations. We also show that people can learn to use semantic relations to make predictions about novel situations, even when they have no prior knowledge of the relations. We also show that people can learn to use semantic relations to make predictions about novel situations, even when they have no prior knowledge of the relations.



these limitations. The first contribution is to introduce and expand the concept learning paradigm. Concept learning studies often focus on individual concepts rather than conceptual systems. This paper presents a game that lends itself well to exploring which concepts and which participants learn best. Participants predict the outcome of representing a concept, transforming one set of numbers into another. Using this game, we find that more concepts are learned more easily by learners and that hard concepts are learned more easily when preceded by a bootstrapping compositional curriculum.

The third contribution is to introduce the idea of using a *meta-language* to guide learning. We propose a computational model of concept learning in which hypotheses represent not merely different definitions of a concept within a fixed LCT, but completely different LCTs. We model learn-

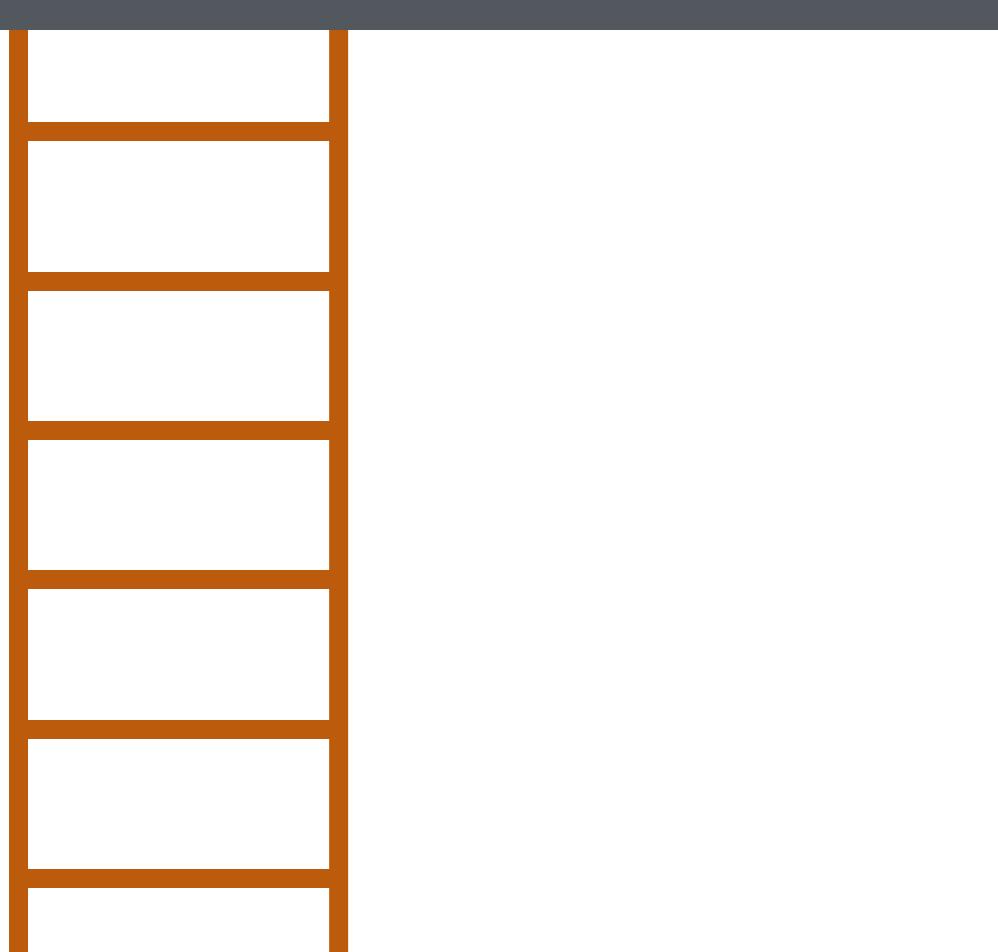


Eric Schulz



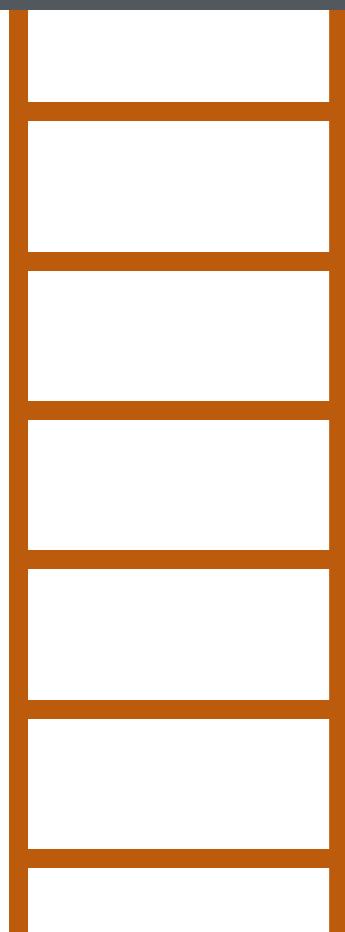
(Rule, Schulz, Piantadosi, Tenenbaum, 2018)

Martha's Magical Machines

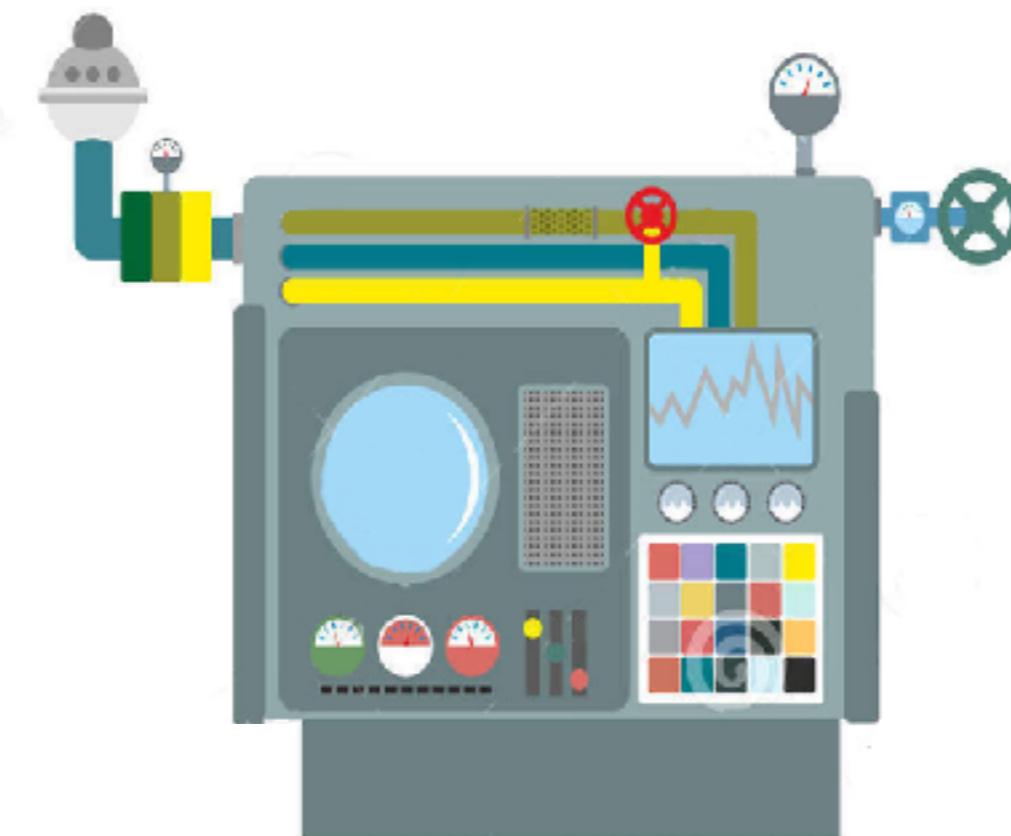
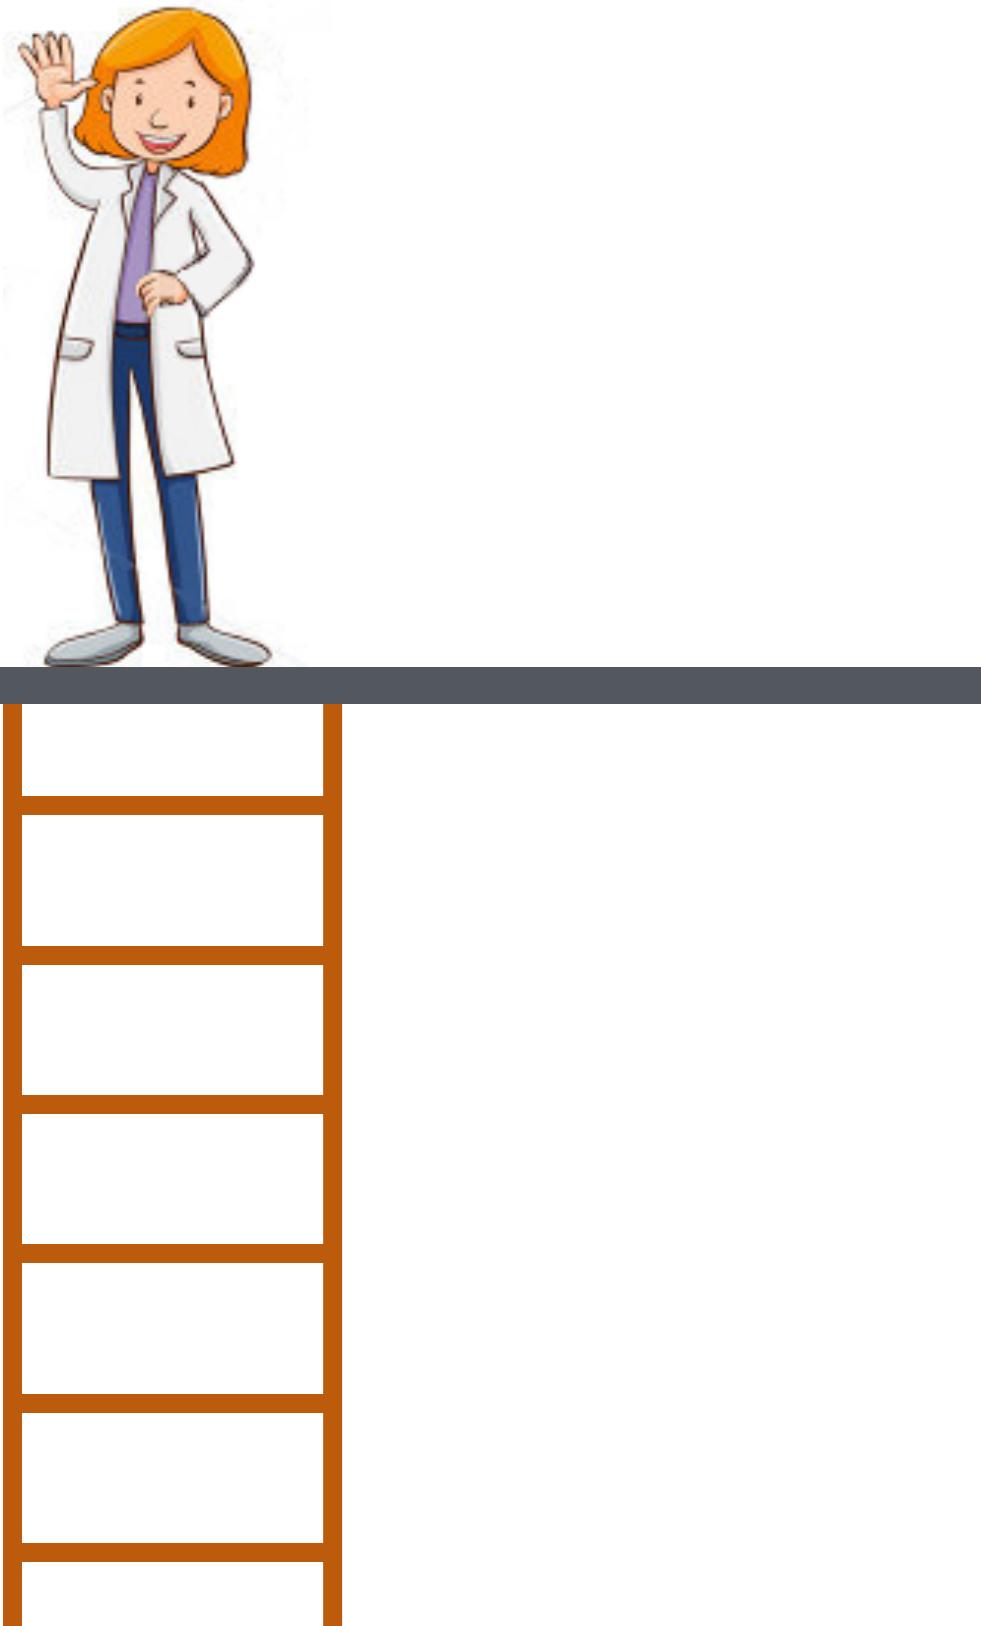


(cf. Piantadosi, Tenenbaum, Goodman, 2016)

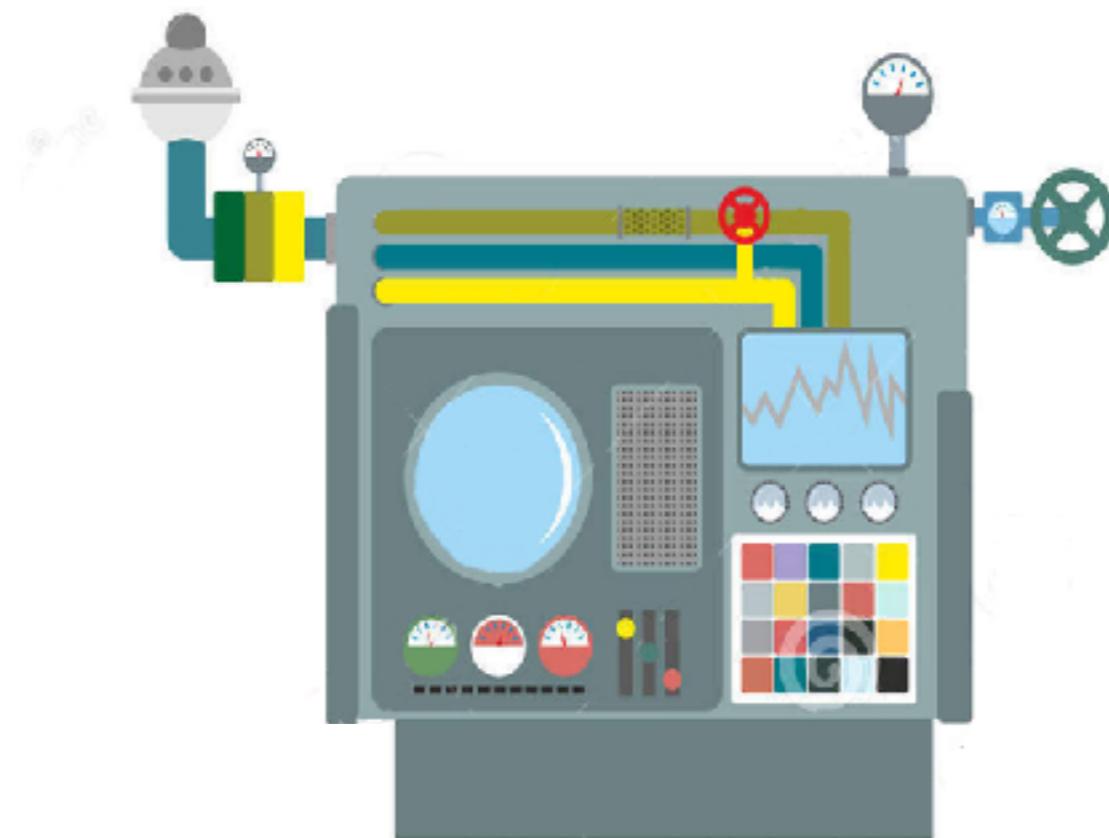
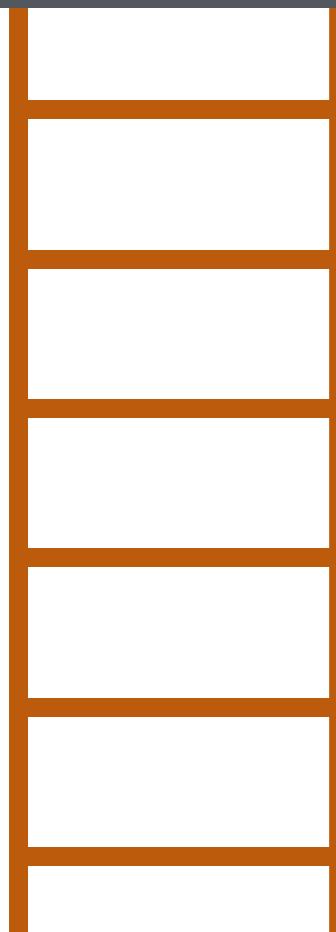
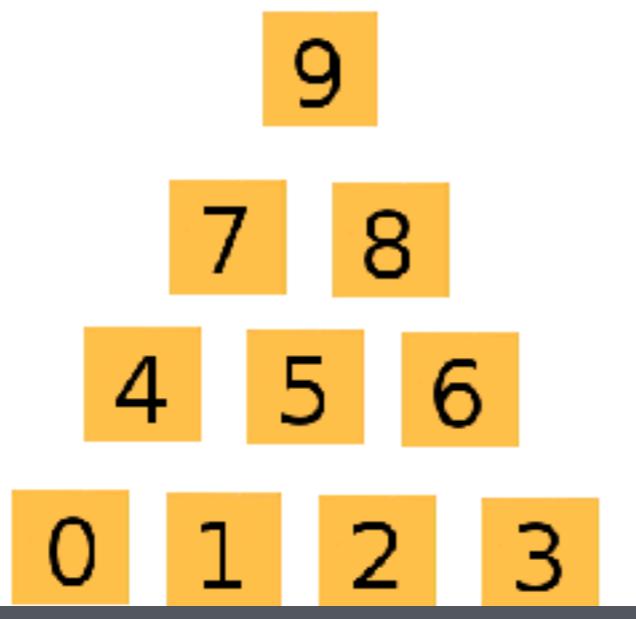
Martha's Magical Machines



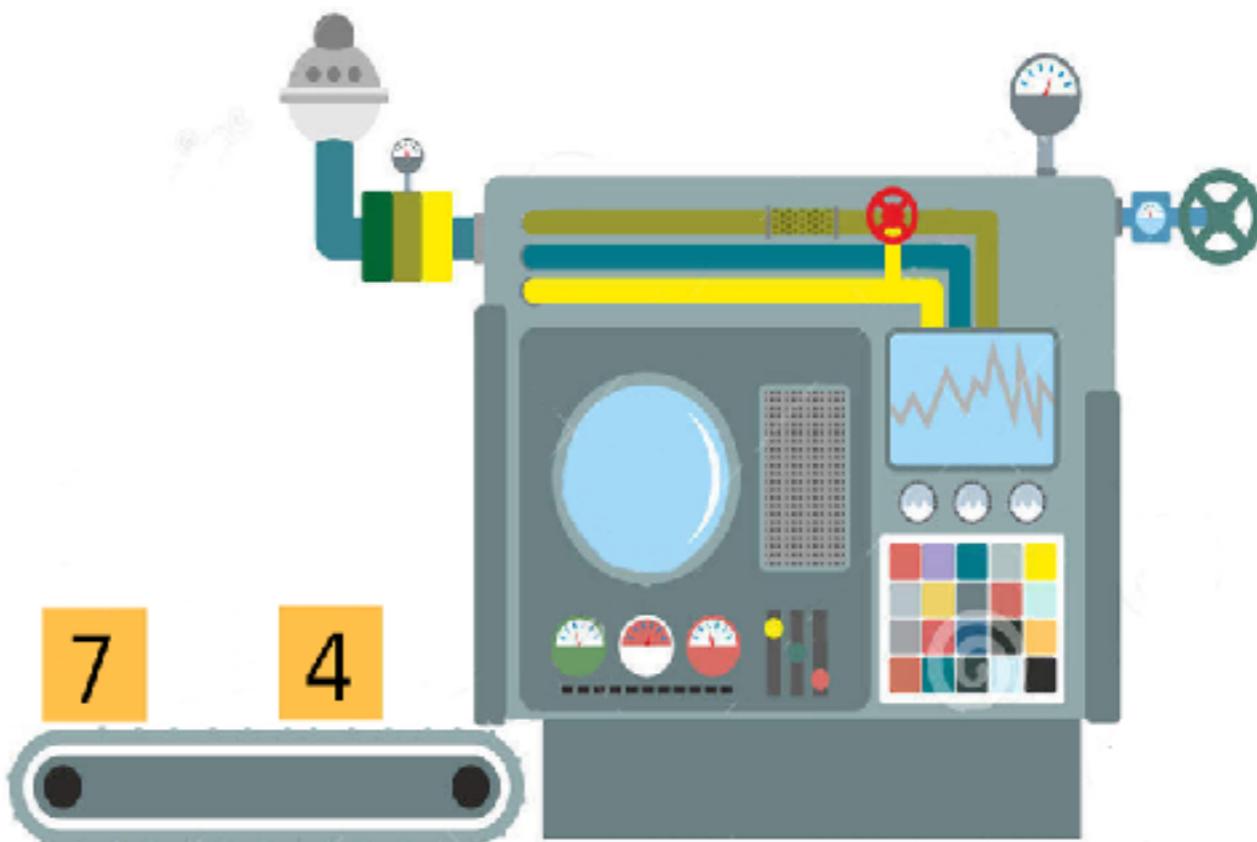
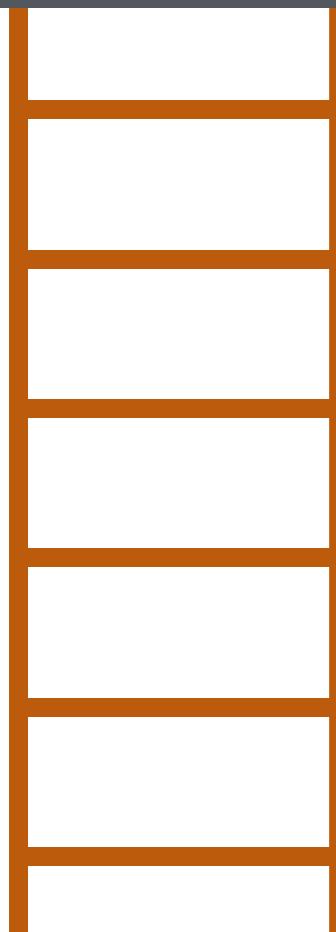
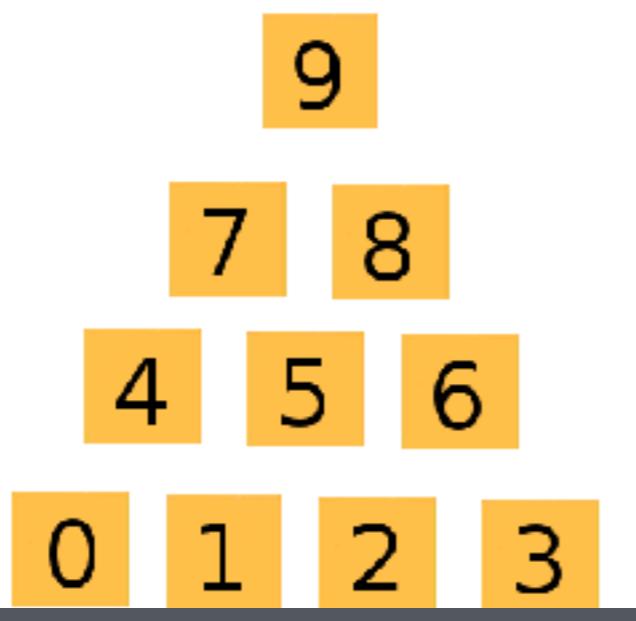
Martha's Magical Machines



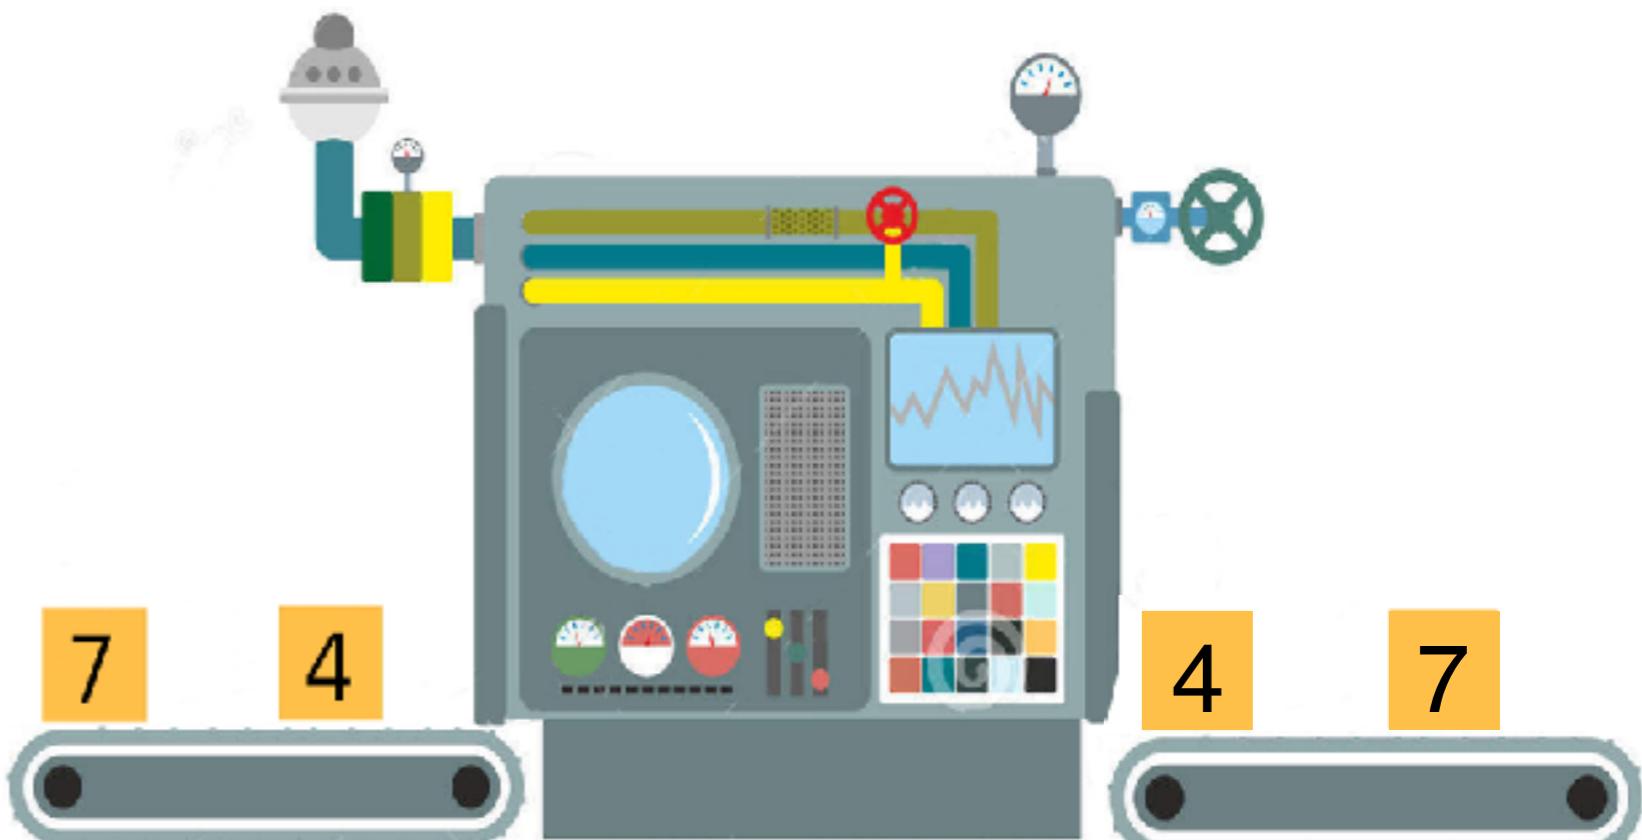
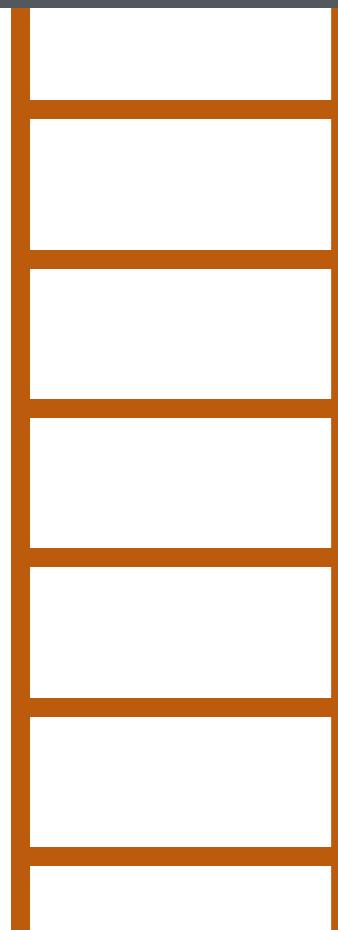
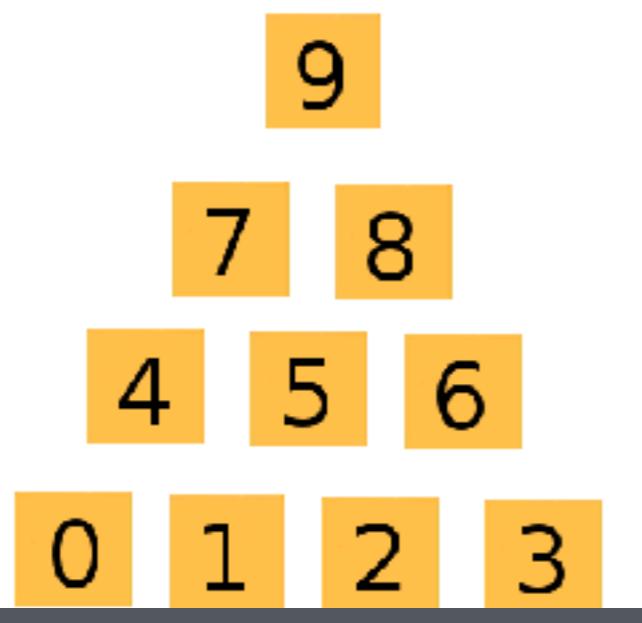
Martha's Magical Machines



Martha's Magical Machines



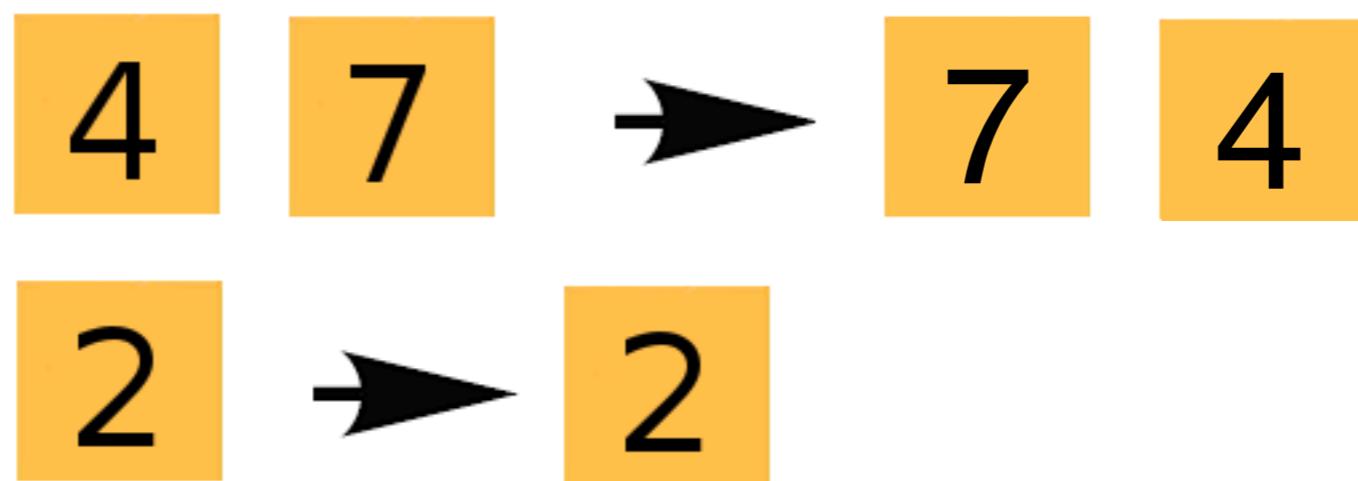
Martha's Magical Machines



Martha's Magical Machines



Martha's Magical Machines



Martha's Magical Machines

4 7 ➤ 7 4

2 ➤ 2

9 6 5 ➤ 5 6 9

Stochastic search over TRSs

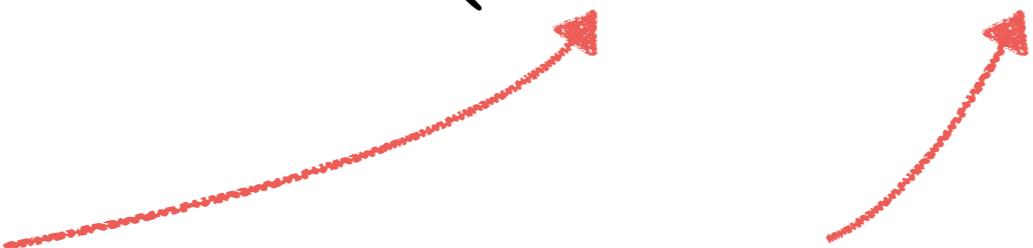
```
def search(data, h0, N=1500, n_top=10, n_steps=50, confidence=2/3):
    dataset = []
    h, score = h0, score(h0)
    hs = heap([(h, score)])
    for (i, o) in data:
        for _ in range(N):
            h_next = propose(h)
            score_next = score(h_next)
            h, score = metropolis(h, score, h_next, score_next)
            hs.insert((h, score))
        best_hs = hs.take_top(n_top)
        o_hat = most_likely_output(i, n_steps, best_hs)
        data.append((i, o))
        N *= (confidence if o_hat == o else 1/confidence)
    return hs
```

Model Primitives

Name & Input/Output Pair	Description
$0, 1, 2$	constant natural numbers
$[]$	the empty list
$\text{succ}(0)$	the successor of x
$\text{cons}(1, [2, 3]) = [1, 2, 3]$	prepend x to y
$\text{sum}([1, 2, 3]) = [6]$	sum x
$\text{add}(3, [1, 2, 3]) = [4, 5, 6]$	add x to the elements of y
$\text{insert}(4, [3, 5]) = [3, 4, 5]$	insert x into y in sorted order
$\text{remove}(1, [6, 1, 4]) = [6, 4]$	remove every x in y
$\text{count}(7, [7, 1, 7]) = [2]$	count every x in y
$\text{even}(5) = \text{false}$	true if x is even else false
$\text{greater}(8, 2) = \text{true}$	true if $x > y$ else false
$\text{if}(\text{true}, [7], [2, 5]) = [7]$	if x then y else z
$\text{nth}(3, [9, 5, 8]) = [8]$	the x^{th} element of y

Model Primitives

$$h_0 = (\Sigma_0^*, R_0)$$



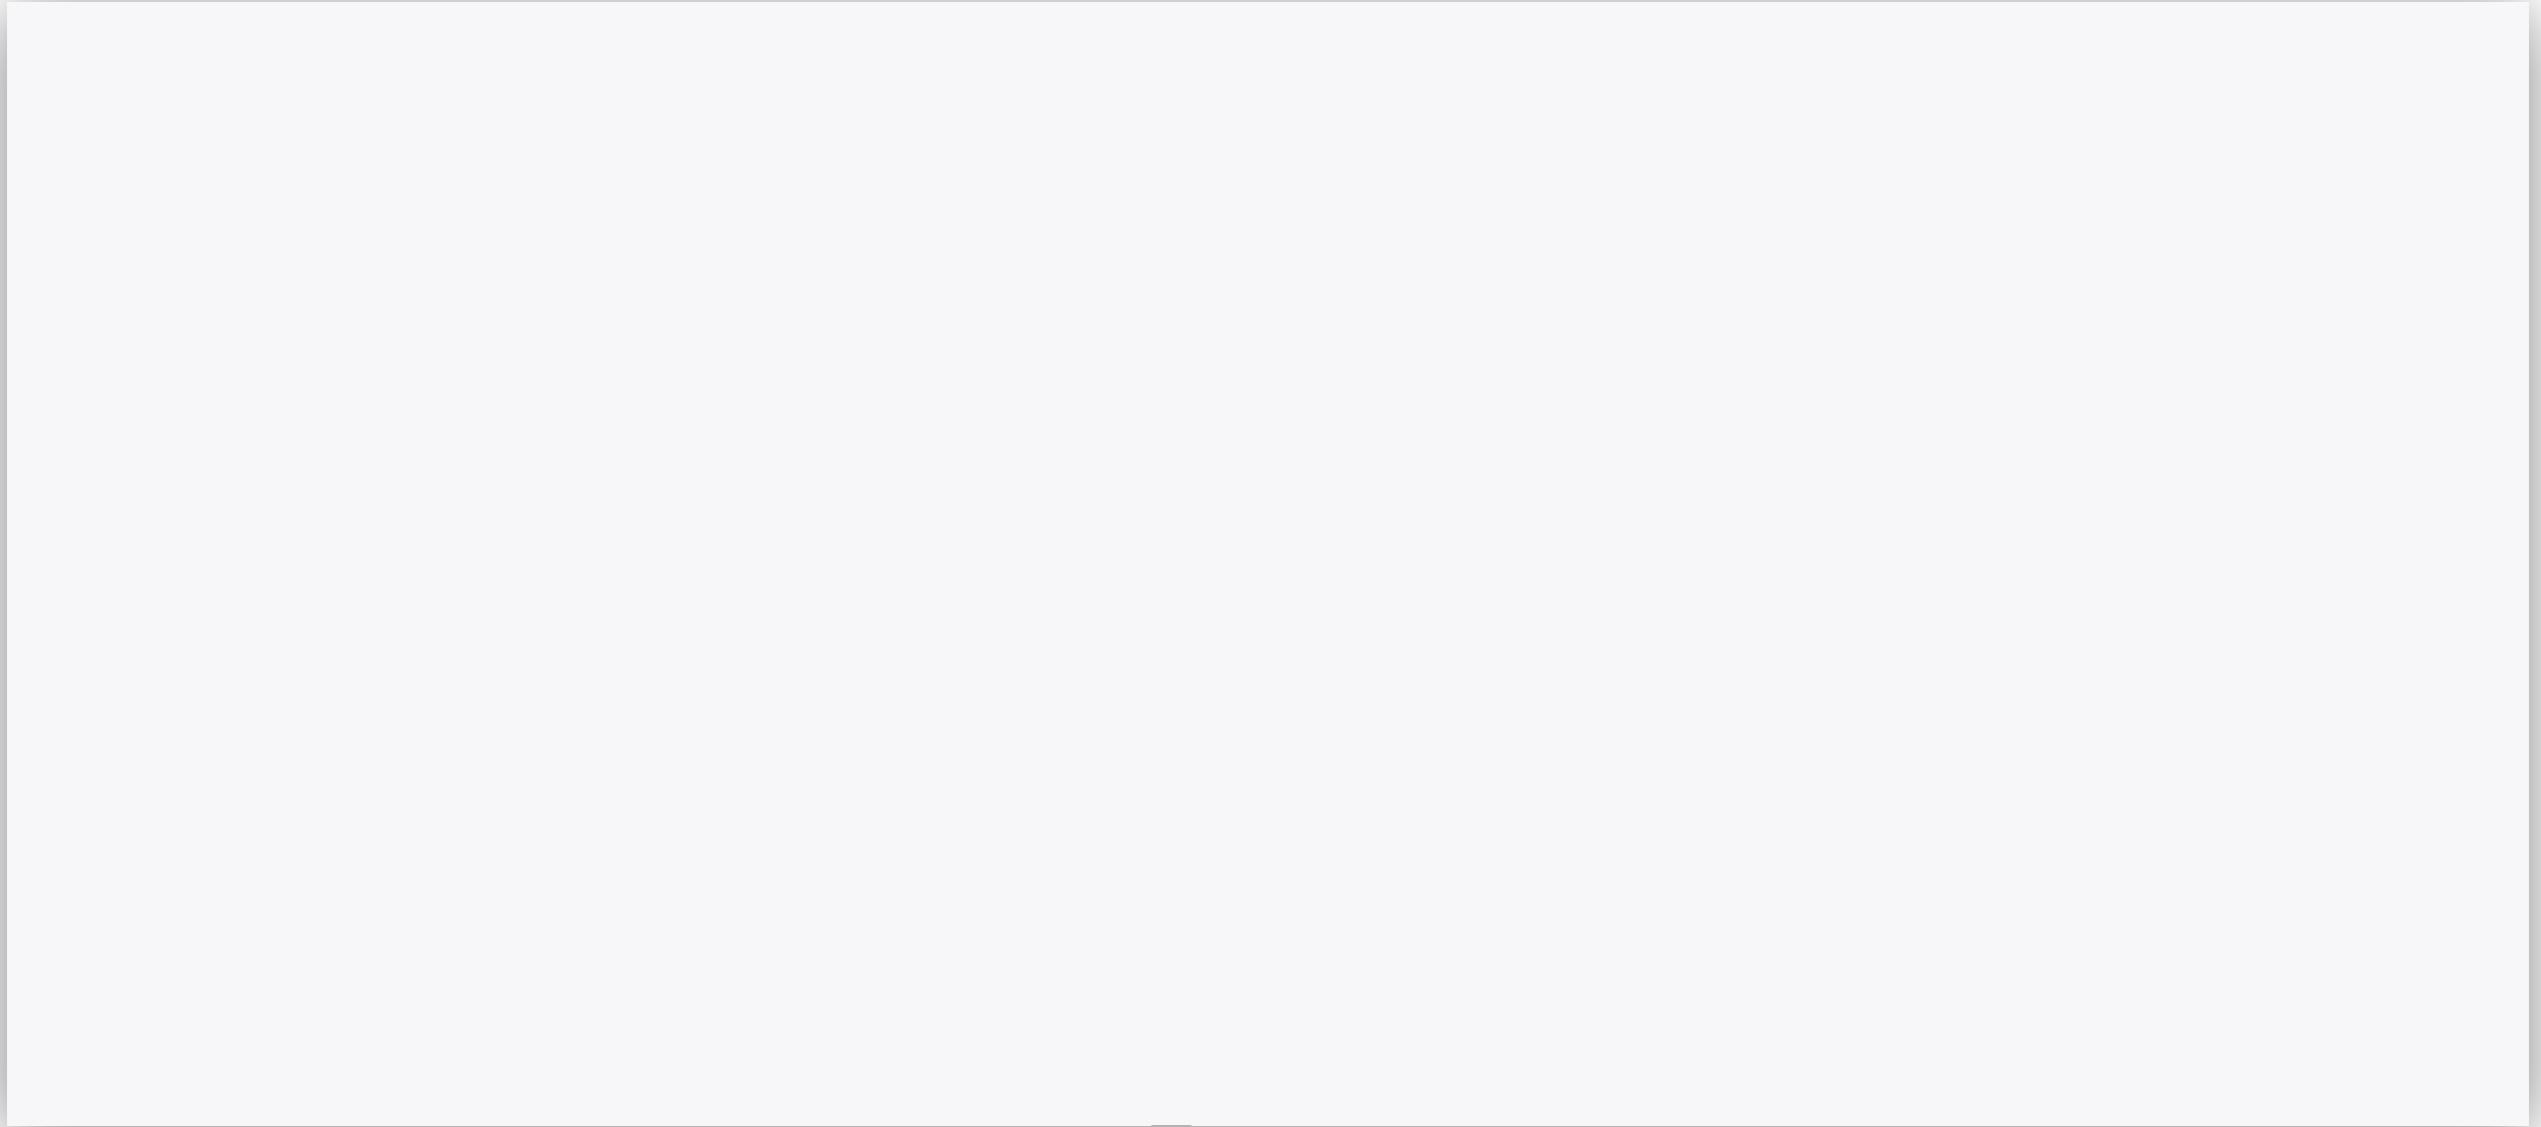
Name & Input/Output Pair	Description
$0, 1, 2$	constant natural numbers
$[]$	the empty list
$\text{succ}(0)$	the successor of x
$\text{cons}(1, [2, 3]) = [1, 2, 3]$	prepend x to y
$\text{sum}([1, 2, 3]) = [6]$	sum x
$\text{add}(3, [1, 2, 3]) = [4, 5, 6]$	add x to the elements of y
$\text{insert}(4, [3, 5]) = [3, 4, 5]$	insert x into y in sorted order
$\text{remove}(1, [6, 1, 4]) = [6, 4]$	remove every x in y
$\text{count}(7, [7, 1, 7]) = [2]$	count every x in y
$\text{even}(5) = \text{false}$	true if x is even else false
$\text{greater}(8, 2) = \text{true}$	true if $x > y$ else false
$\text{if}(\text{true}, [7], [2, 5]) = [7]$	if x then y else z
$\text{nth}(3, [9, 5, 8]) = [8]$	the x^{th} element of y

*plus the target concept

Experiment 1

- ▶ **149 participants** (61 female, mean age=36.93, SD=12.20)
- ▶ **5 concepts/participant** (out of 12)
- ▶ **10 trials/concept**

Experiment 1



Experiment 1

```
# const xs: return 3
# Example: const([1,2,4]) = [3]
const(x_) = 3;
```

Experiment 1

```
# const xs: return 3
# Example: const([1,2,4]) = [3]
const(x_) = 3;
```

```
# index-in-head xs: return the headth element of the xs
# Example: index_in_head([2,3]) = [3]
index-in-head(cons(0 y_)) = 0
index-in-head(cons(succ(x_) y_)) = nth(x_ y_);
```

Experiment 1

```
# const xs: return 3
# Example: const([1,2,4]) = [3]
const(x_) = 3;

# total xs: sum all the elements of xs
# Example: total([1,2,3]) = [6]
total(x_) = sum(x_);

# increment xs: add 1 to each element of xs
# Example: increment([1,2]) = [2,3]
increment(x_) = add(1 x_);

# head xs: return the first element of xs
# Example: head([2,3,1]) = [2]
head(cons(x_ y_)) = x_;

# length xs: compute the length of xs
# Example: length([2,3,1]) = [3]
length([]) = 0;
length(cons(x_ y_)) = succ(length(y_));

# sort xs: sort xs
# Example: sort([3,1]) = [1,3]
sort([]) = [];
sort(cons(x_ y_)) = insert(x_ sort(y_));

# deduplicate xs: remove all duplicates from xs
# Example: deduplicate([2,1,2,2,1]) = [2,1]
deduplicate([]) = [];
deduplicate(cons(x_ y_)) =
    cons(x_ deduplicate(remove(x_ y_))))
```

```
# cumsum xs: cumulatively sum the elements of xs
# Example: cumsum([2,3,1]) = [2,5,6]
cumsum([]) = [];
cumsum(cons(x_ y_)) = cons(x_ cumsum(add(x_ y_)));

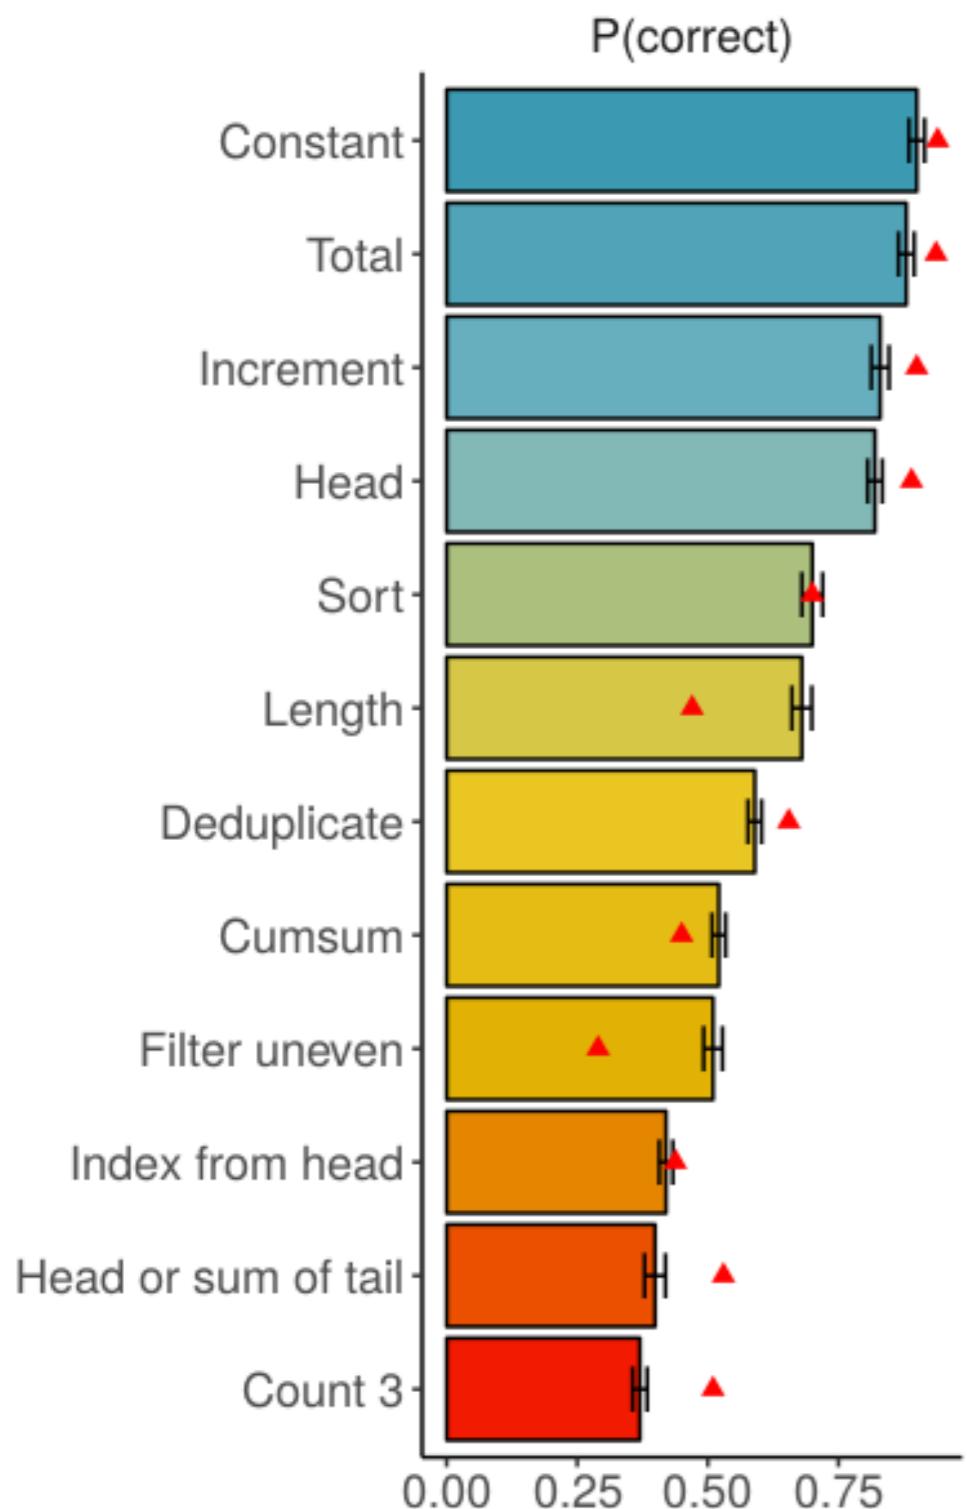
# filter_odd xs: remove the odd numbers from xs
# Example: filter_odd([2,3,1,4]) = [2,4]
filter_odd([]) = [];
filter_odd(cons(x_ y_)) =
    if(even?(x_)) cons(x_ filter_odd(y_)) filter_odd(y_);

# index-in-head xs: return the headth element of the xs
# Example: index_in_head([2,3]) = [3]
index-in-head(cons(0 y_)) = 0
index-in-head(cons(succ(x_) y_)) = nth(x_ y_);

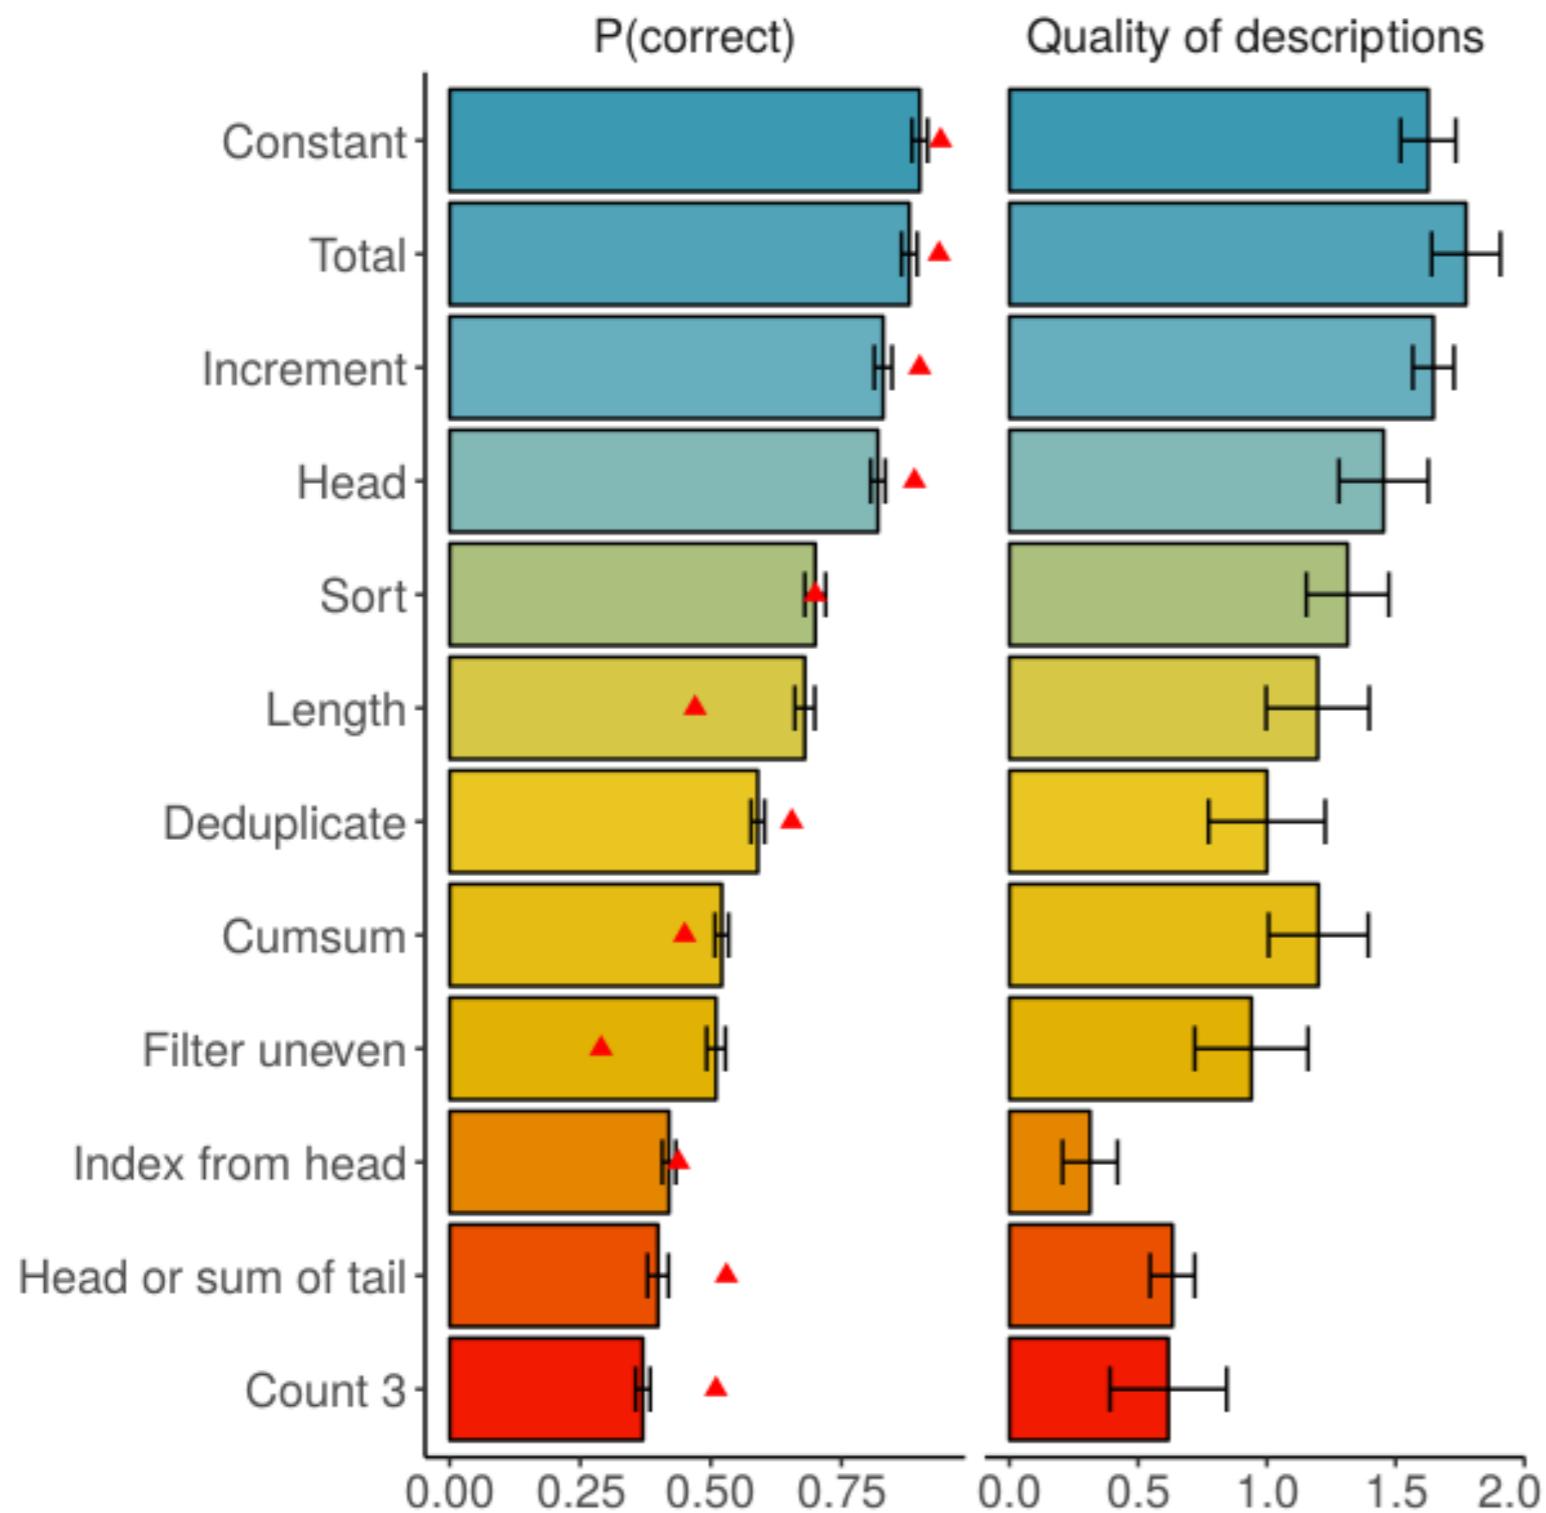
# head-or-tail: return the larger of head or sum-of-tail
# Example: head_or_tail([2,3,1]) = [4]
head-or-tail([]) = 0;
head-or-tail(cons(x_ y_)) =
    if(greater(x_ sum(y_)) x_ sum(y_));

# count3 xs: how often does 3 appear in xs?
# Example: count3([2,3,3]) = [2]
count3(x_) = count(succ(succ(succ(0))) x_);
```

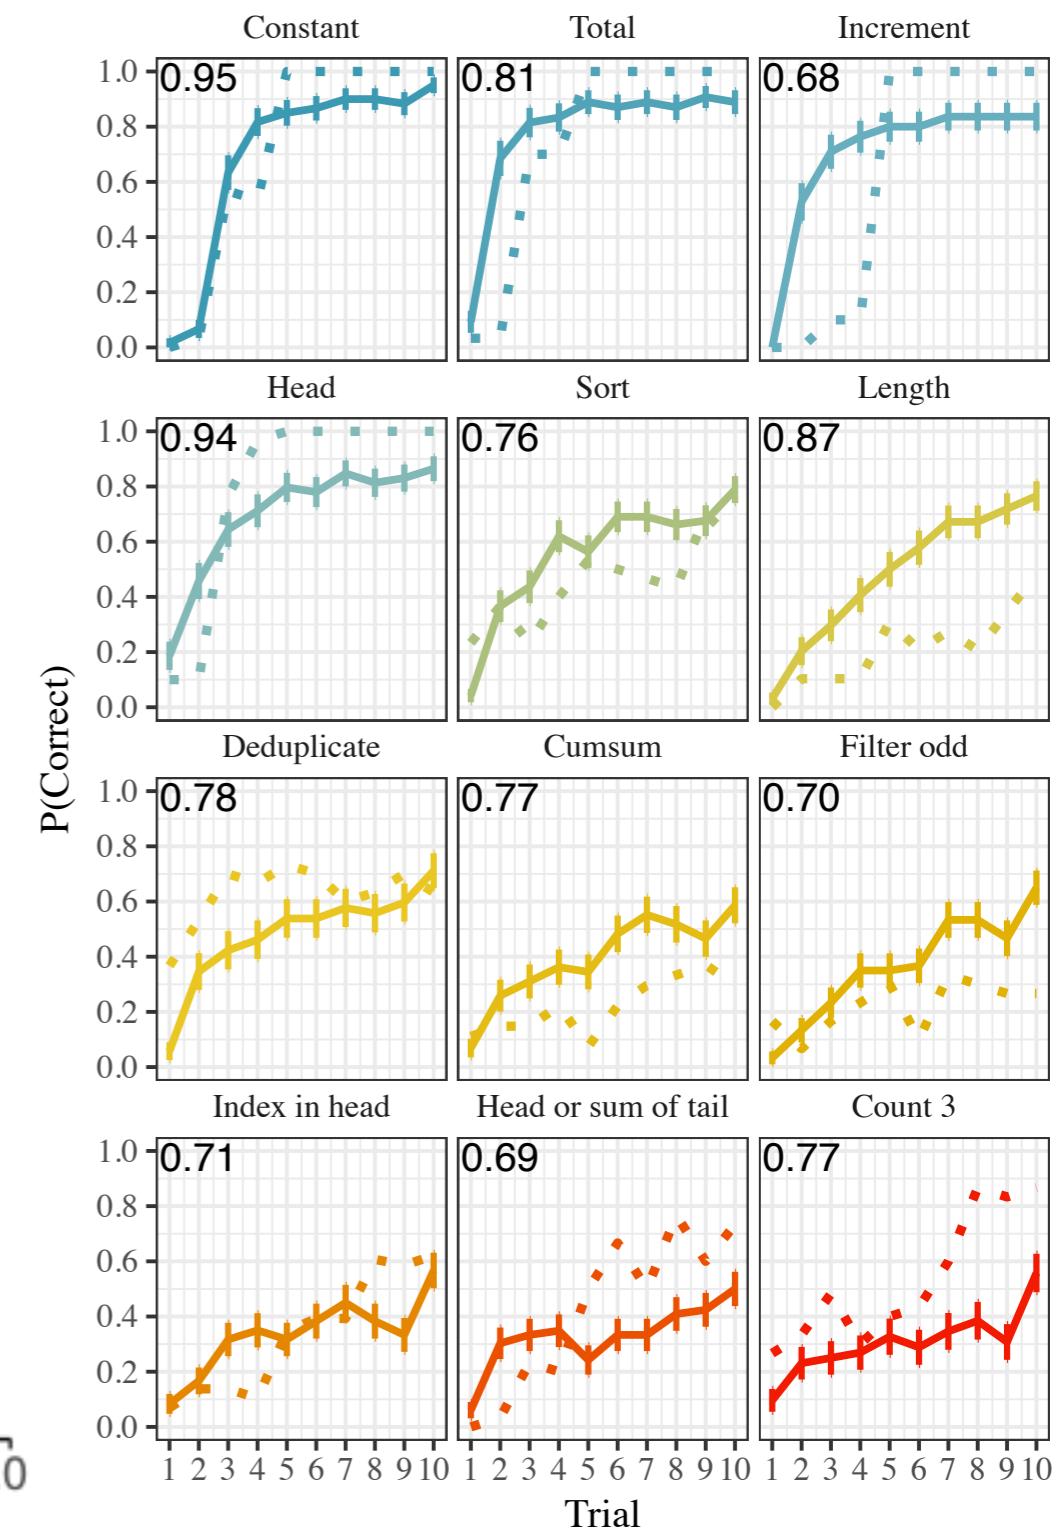
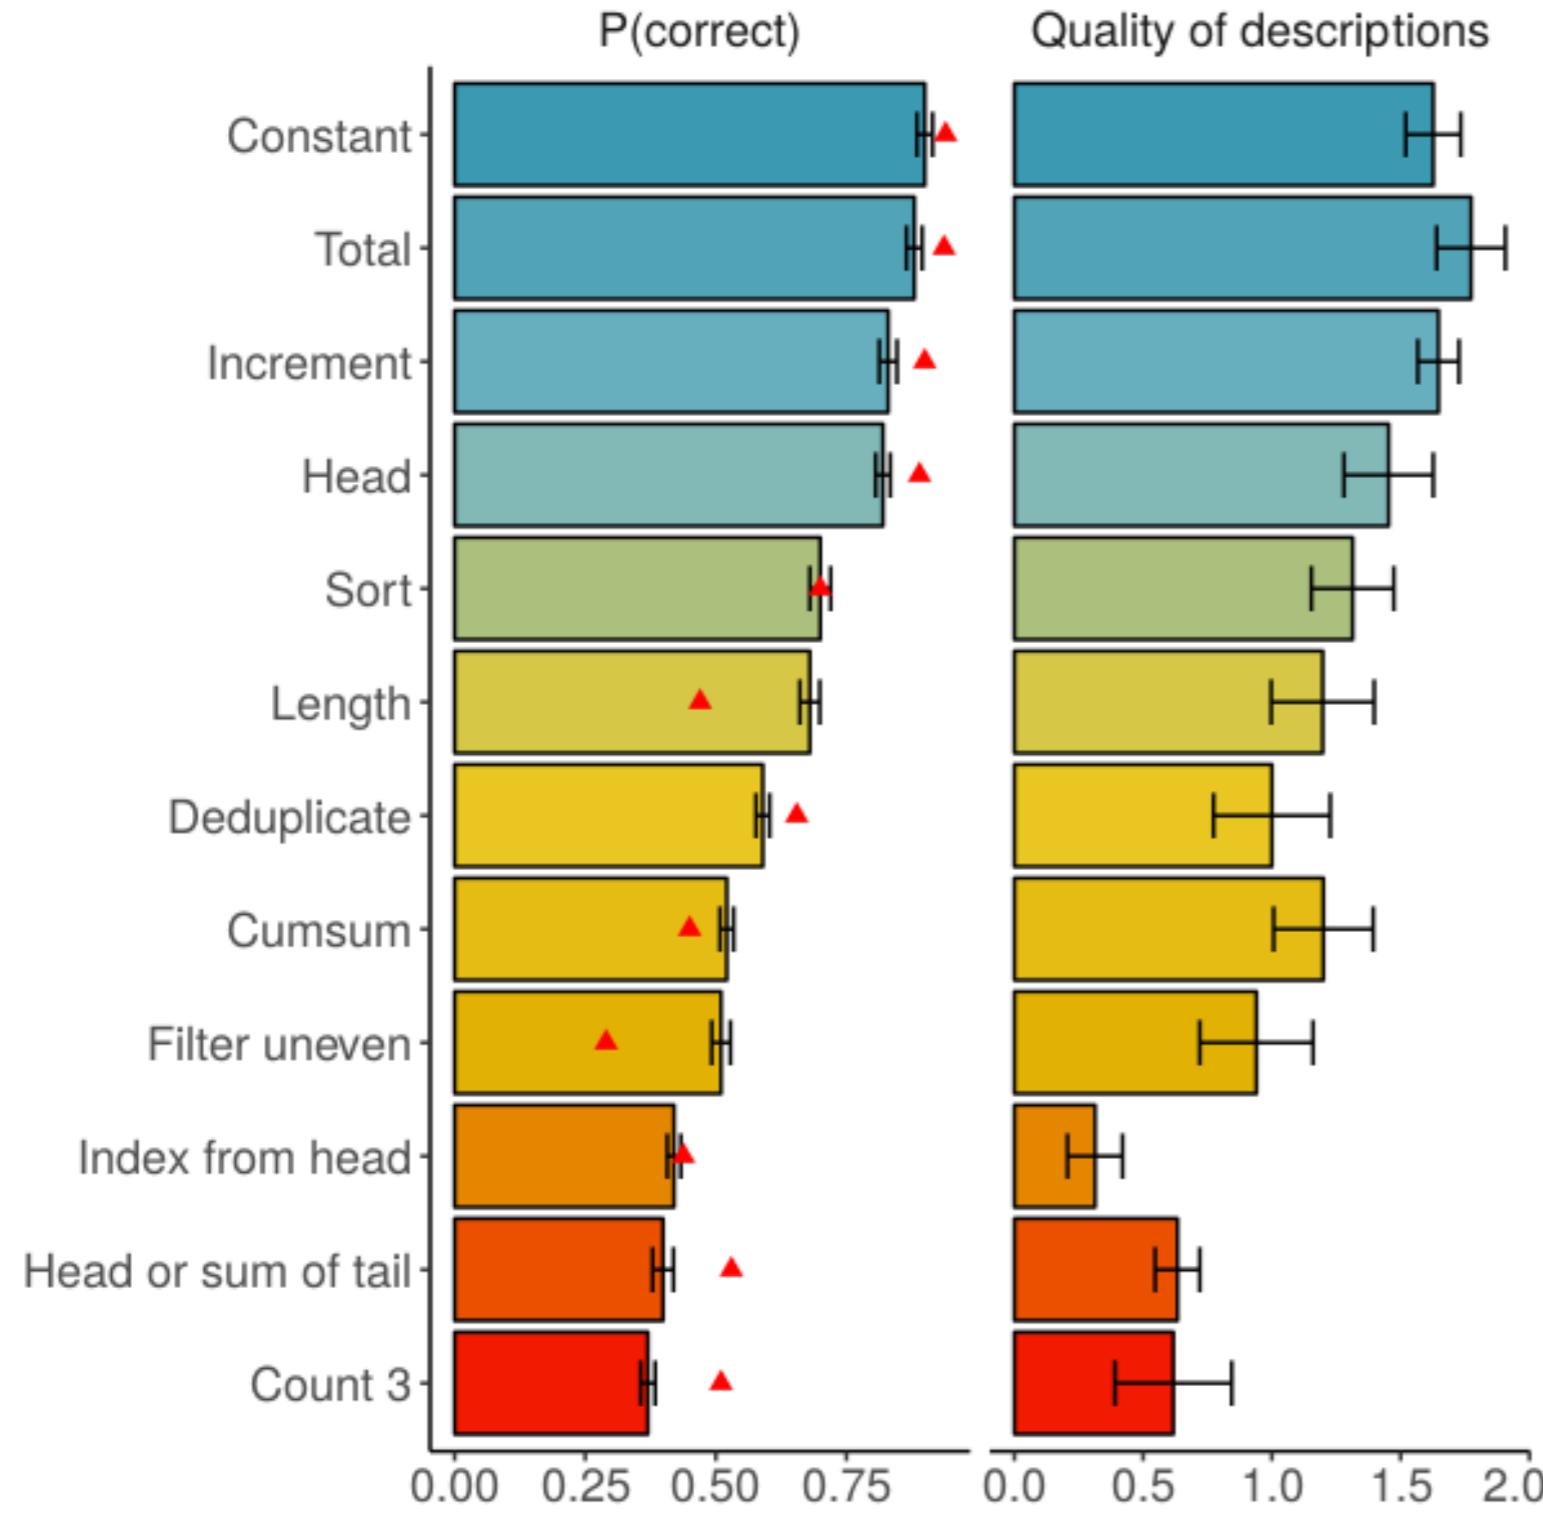
Experiment 1



Experiment 1



Experiment 1



Experiment 2

- ▶ **91 participants** (46 males, mean age=34.51, SD=10.57)
- ▶ randomly assigned condition:
relevant curriculum or random curriculum
- ▶ 4 concepts/condition
- ▶ 10 trials/concept

Experiment 2

```
# head xs: return the first element of xs          # count3 xs: how often does 3 appear in xs?  
# Example: head([2,3,1]) = [2]                    # Example: count3([3,2,3]) = [2]  
head(cons(x_ y_)) = x_;                         count3(x_) = count(succ(succ(succ(0))) x_);  
  
# tail xs: return all but the first element of xs # count-head-in-tail xs: how often is head in the tail?  
# Example: tail([2,3,3]) = [3,3]                  # Example: count-head-in-tail([2,3,2]) = [1]  
tail([]) = [];                                    count-head-in-tail([]) = 0;  
tail(cons(x_ y_)) = y_;                         count-head-in-tail(x_) = count(head(x_) tail(x_));
```

Experiment 2

```
# head xs: return the first element of xs
# Example: head([2,3,1]) = [2]
head(cons(x_ y_)) = x_;

# tail xs: return all but the first element of xs
# Example: tail([2,3,3]) = [3,3]
tail([]) = [];
tail(cons(x_ y_)) = y_;

# count3 xs: how often does 3 appear in xs?
# Example: count3([3,2,3]) = [2]
count3(x_) = count(succ(succ(succ(0))) x_);

# count-head-in-tail xs: how often is head in the tail?
# Example: count-head-in-tail([2,3,2]) = [1]
count-head-in-tail([]) = 0;
count-head-in-tail(x_) = count(head(x_) tail(x_));

# const xs: return 3
# Example: const([1,2,4]) = [3]
const(x_) = 3;

# total xs: sum all the elements of xs
# Example: total([1,2,3]) = [6]
total(x_) = sum(x_);

# increment xs: add 1 to each element of xs
# Example: increment([1,2]) = [2,3]
increment(x_) = add(1 x_);

# length xs: compute the length of xs
# Example: length([2,3,1]) = [3]
length([]) = 0;
length(cons(x_ y_)) = succ(length(y_));

# sort xs: sort xs
# Example: sort([3,1]) = [1,3]
sort([]) = [];
sort(cons(x_ y_)) = insert(x_ sort(y_));

# deduplicate xs: remove all duplicates from xs
# Example: deduplicate([2,1,2,2,1]) = [2,1]
deduplicate([]) = [];
deduplicate(cons(x_ y_)) =
    cons(x_ deduplicate(remove(x_ y_)))



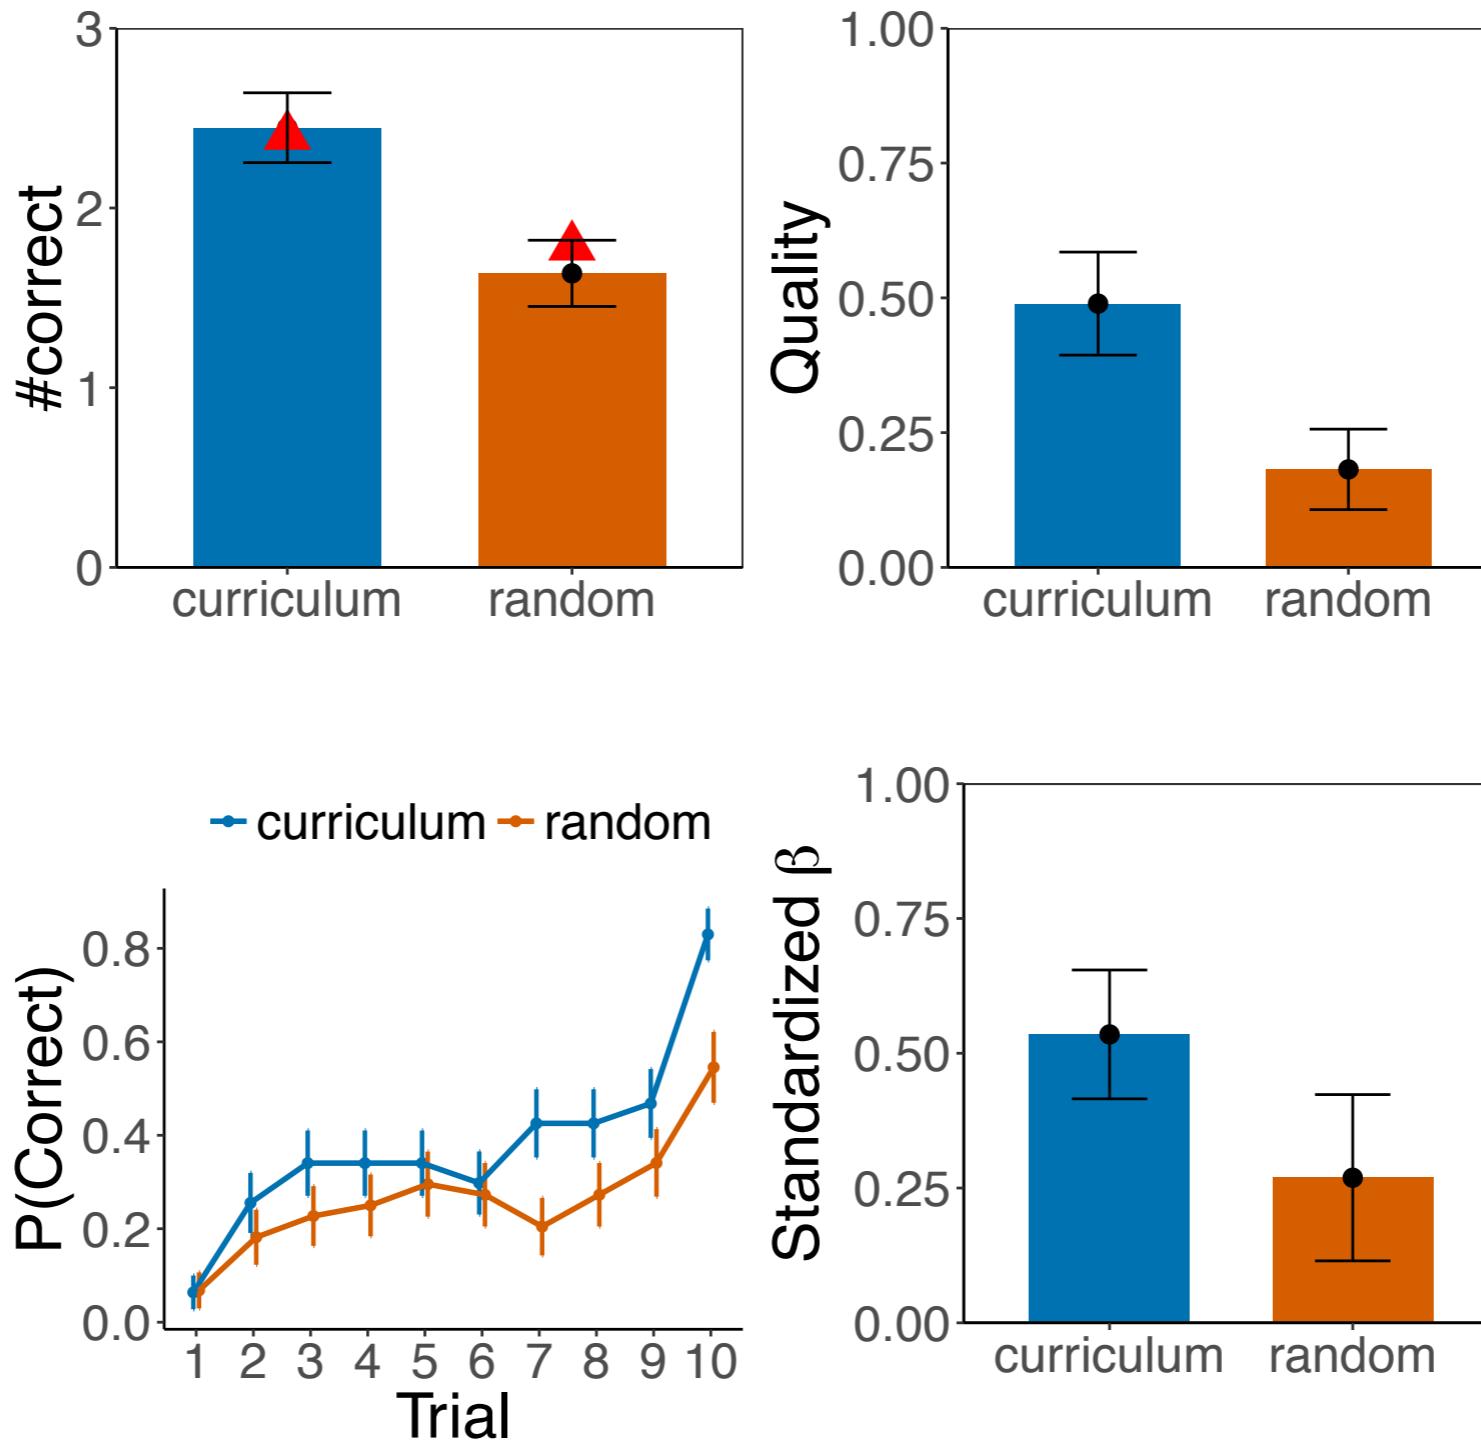
# cumsum xs: cumulatively sum the elements of xs
# Example: cumsum([2,3,1]) = [2,5,6]
cumsum([]) = [];
cumsum(cons(x_ y_)) = cons(x_ cumsum(add(x_ y_)));

# filter_odd xs: remove the odd numbers from xs
# Example: filter_odd([2,3,1,4]) = [2,4]
filter_odd([]) = [];
filter_odd(cons(x_ y_)) =
    if(even?(x_) cons(x_ filter_odd(y_)) filter_odd(y_));

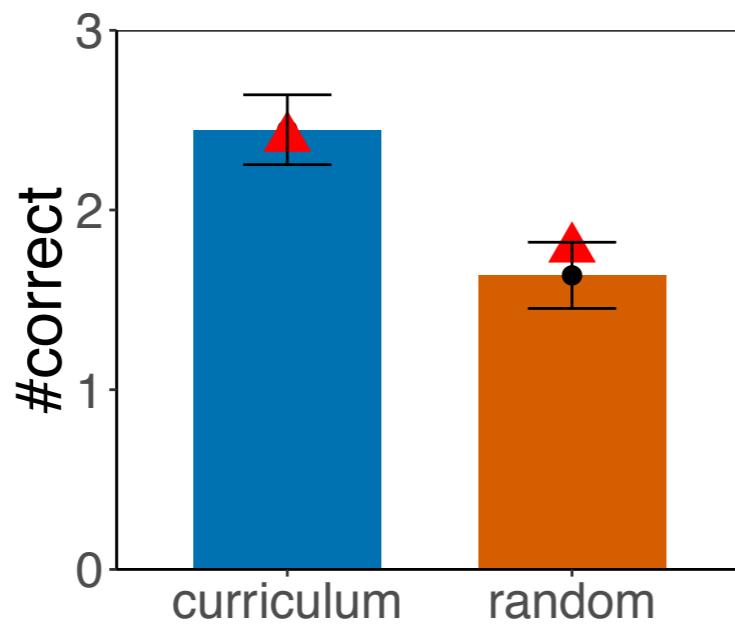
# index-in-head xs: return the headth element of the xs
# Example: index_in_head([2,3]) = [3]
index-in-head(cons(0 y_)) = 0;
index-in-head(cons(succ(x_) y_)) = nth(x_ y_);

# head-or-tail: return the larger of head or sum-of-tail
# Example: head_or_tail([2,3,1]) = [4]
head-or-tail([]) = 0;
head-or-tail(cons(x_ y_)) =
    if(greater(x_ sum(y_)) x_ sum(y_));
```

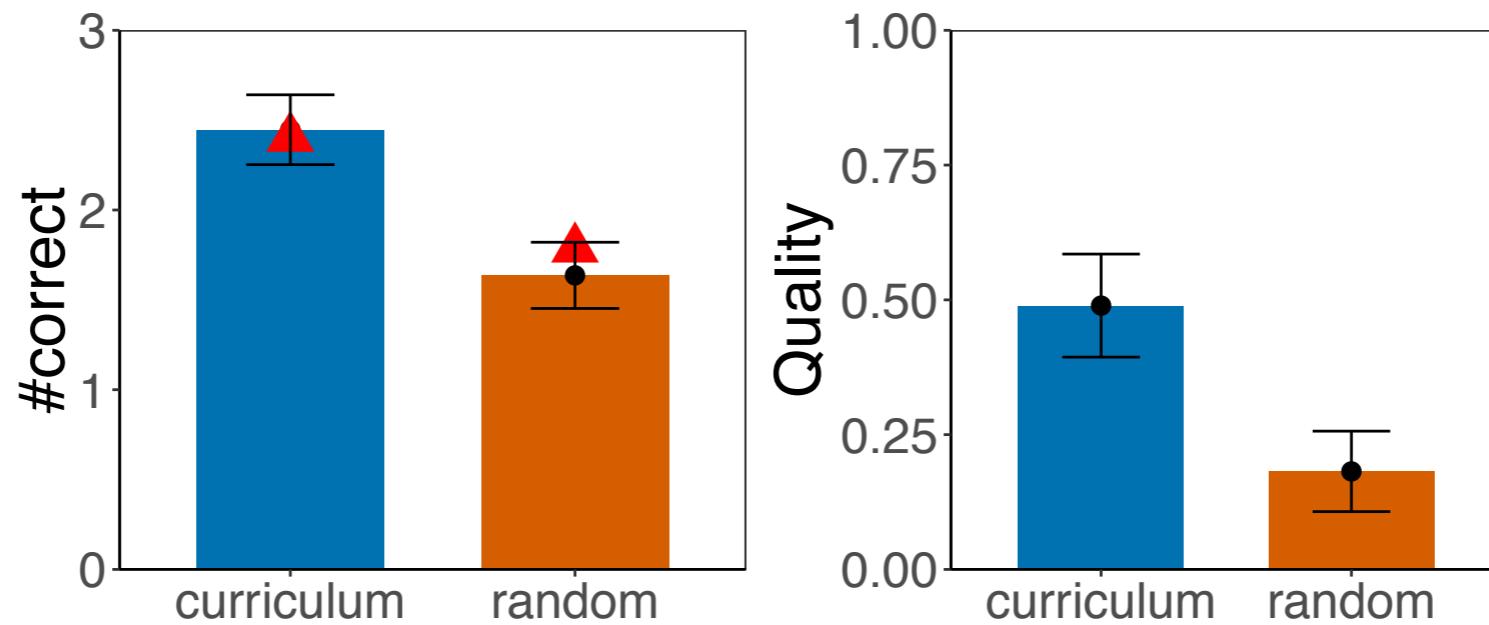
Experiment 2



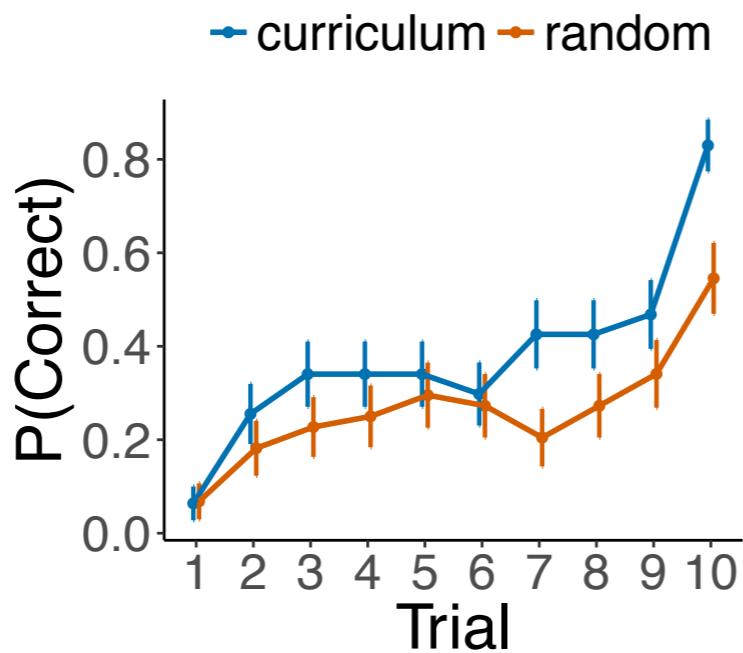
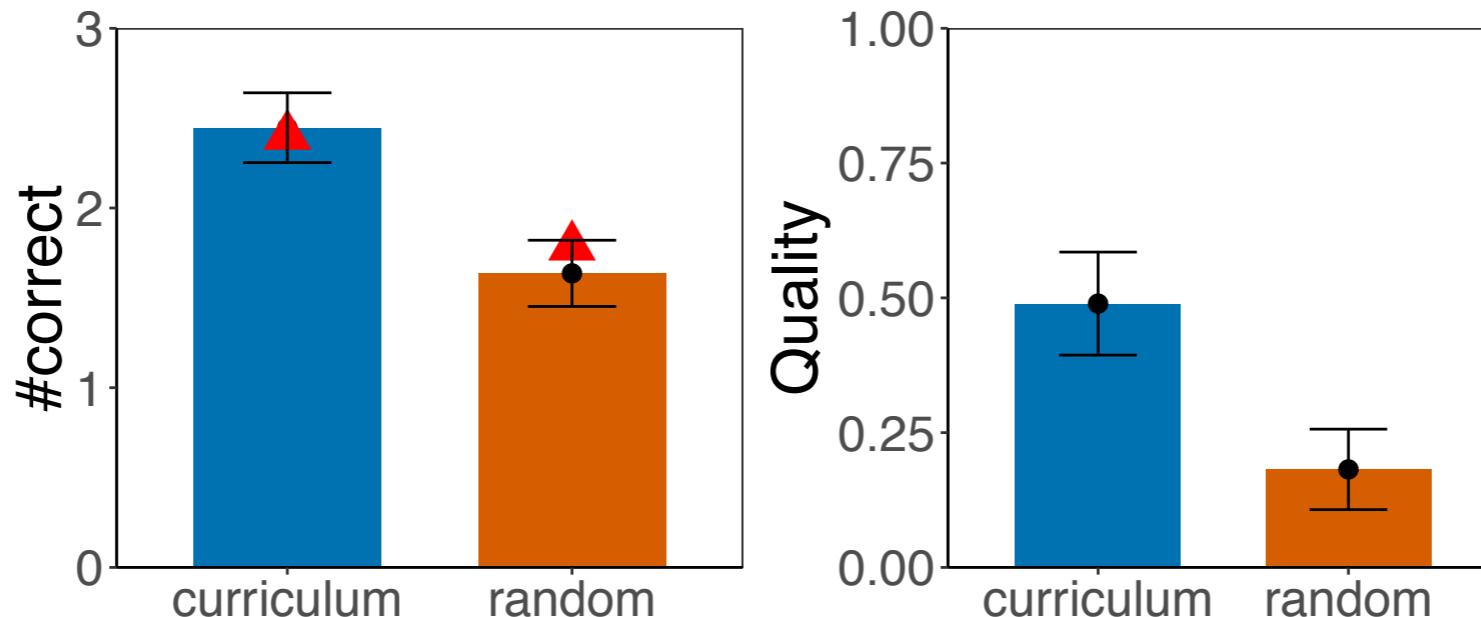
Experiment 2



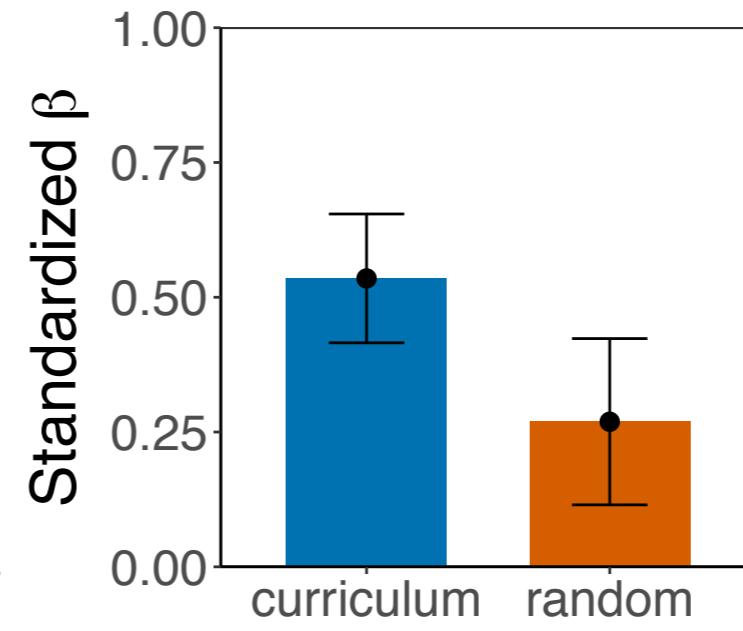
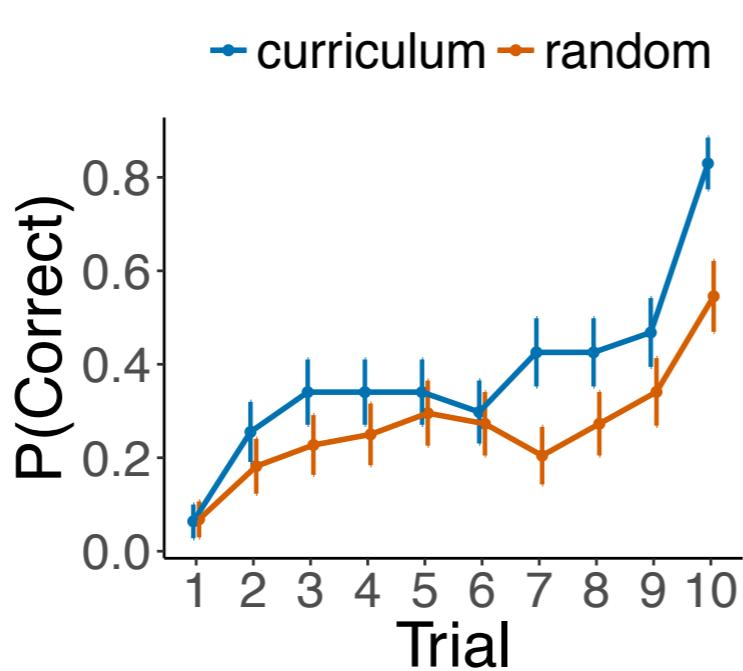
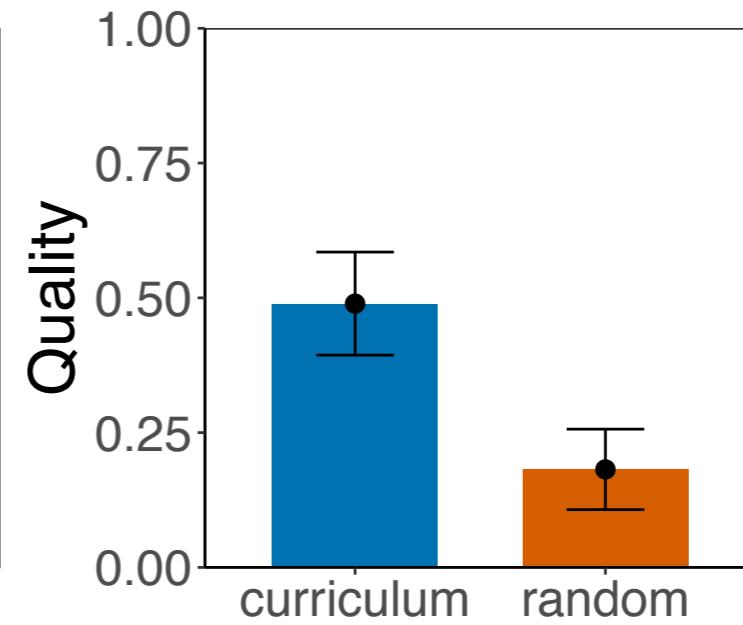
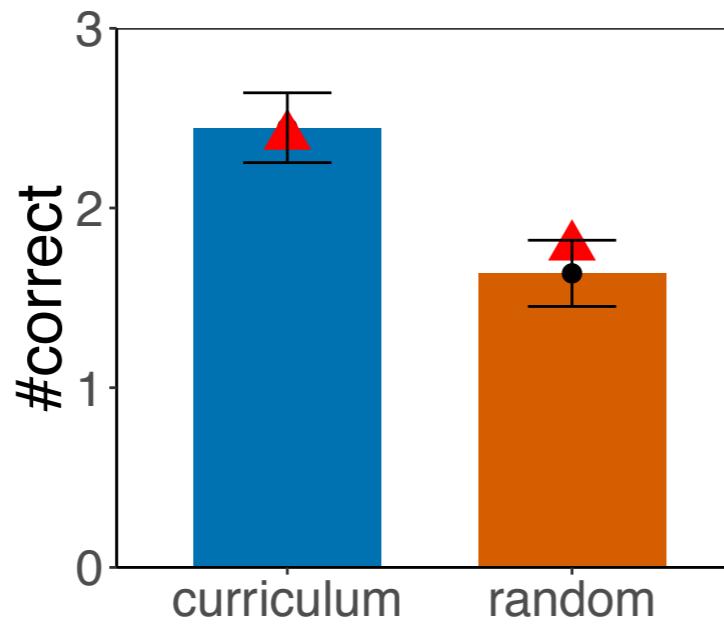
Experiment 2



Experiment 2



Experiment 2



This talk

- ▶ learning as programming
- ▶ bootstrapping the LOT with term rewriting
- ▶ toward a model of conceptual change

This talk

- ▶ learning as programming
- ▶ bootstrapping the LOT with term rewriting
- ▶ toward a model of conceptual change

Thank you!

